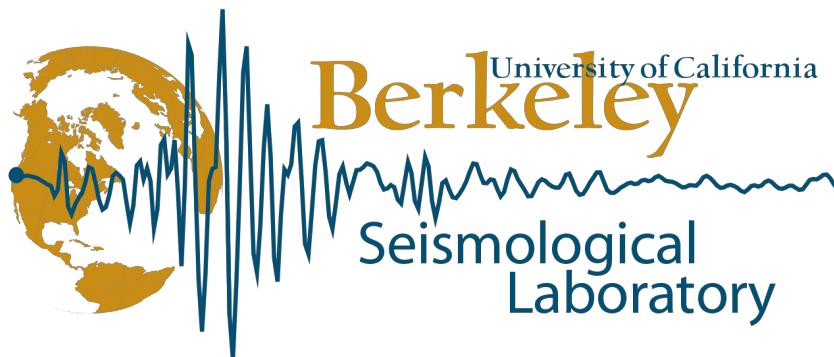




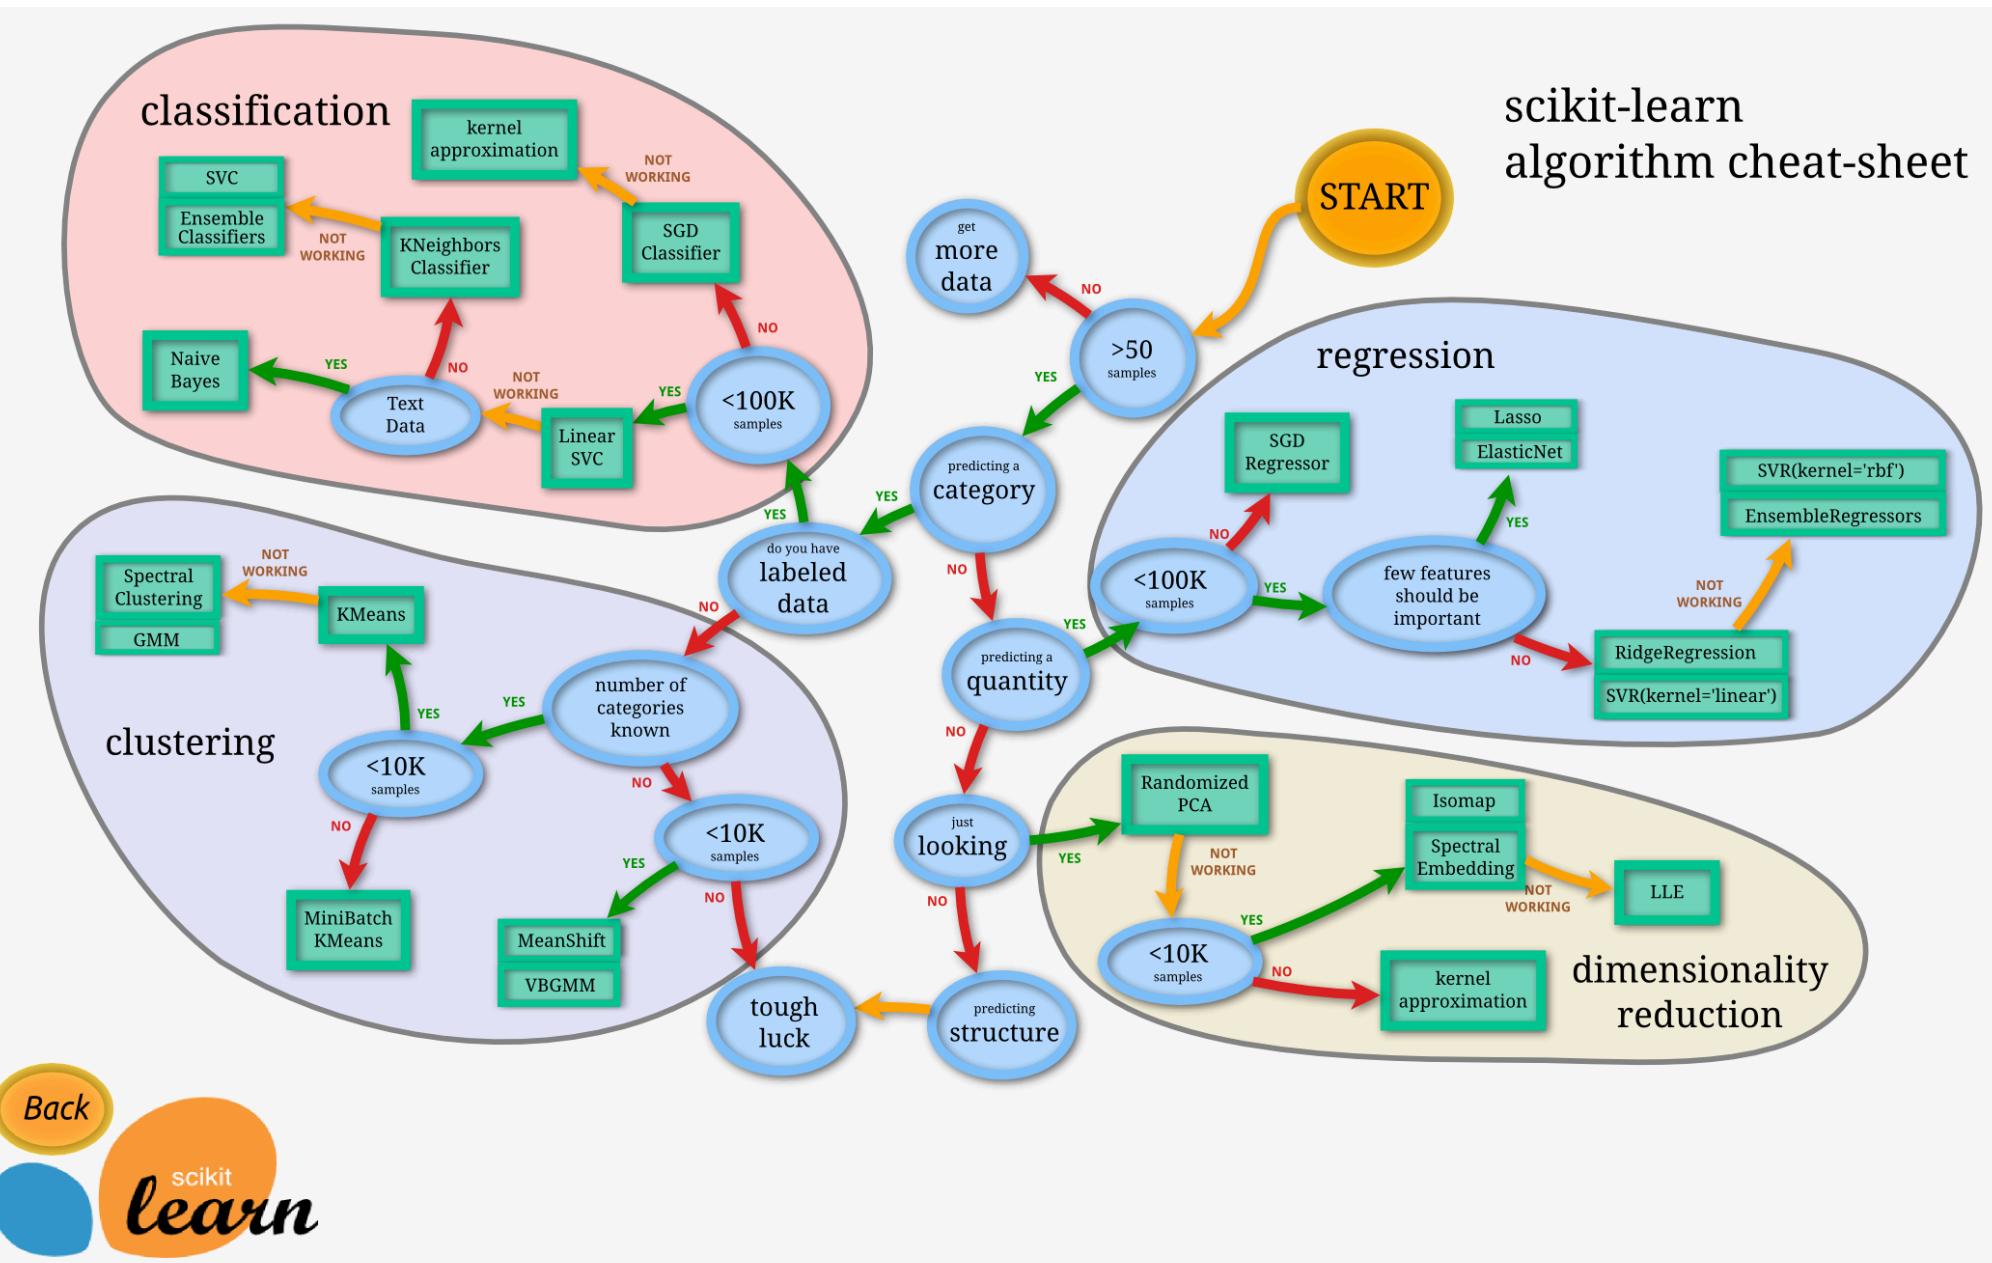
Classification

Qingkai Kong



<http://seismo.berkeley.edu/qingkaikong/>

scikit-learn algorithm cheat-sheet



Interactive version: http://scikit-learn.org/stable/tutorial/machine_learning_map/index.html



Pipeline of training a machine learning model.

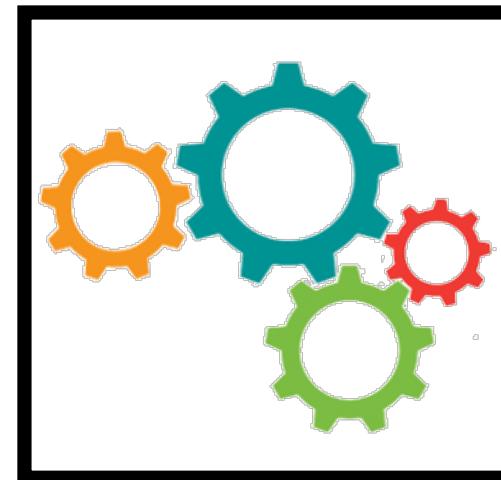
Data examples

01100
10110
11110

Optimization algorithm



Tunable Model



Trained Model

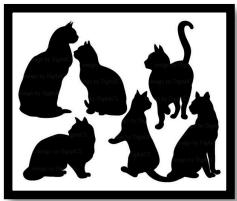
MAKE
THINGS
HAPPEN!

Representation of data

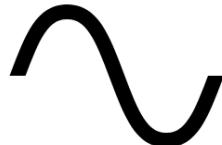
Raw data



Documents



Images



Numbers

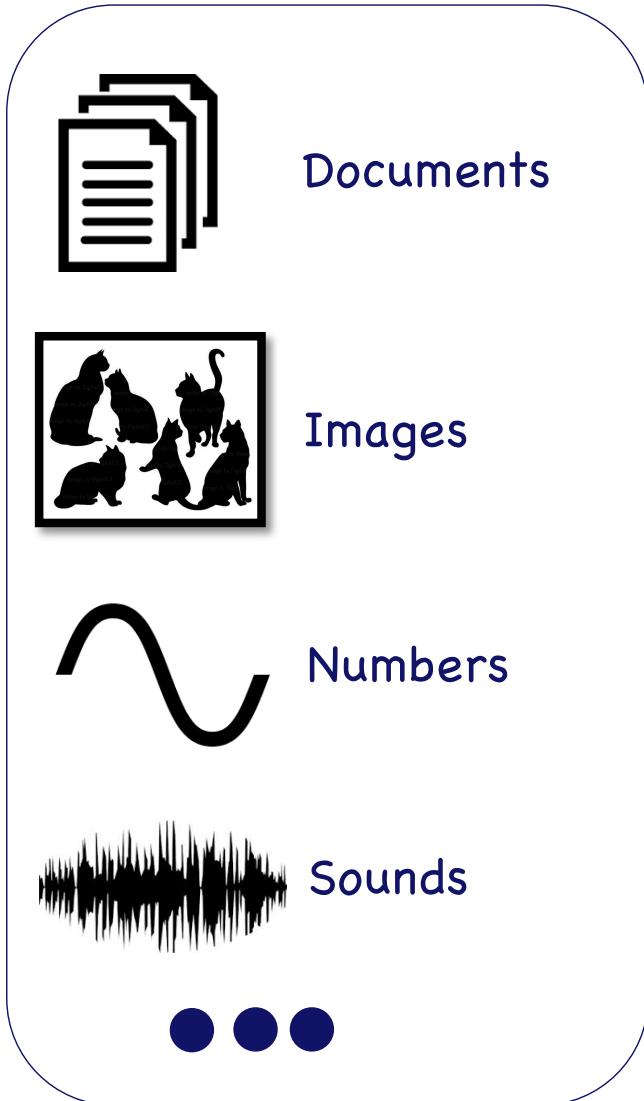


Sounds

• • •

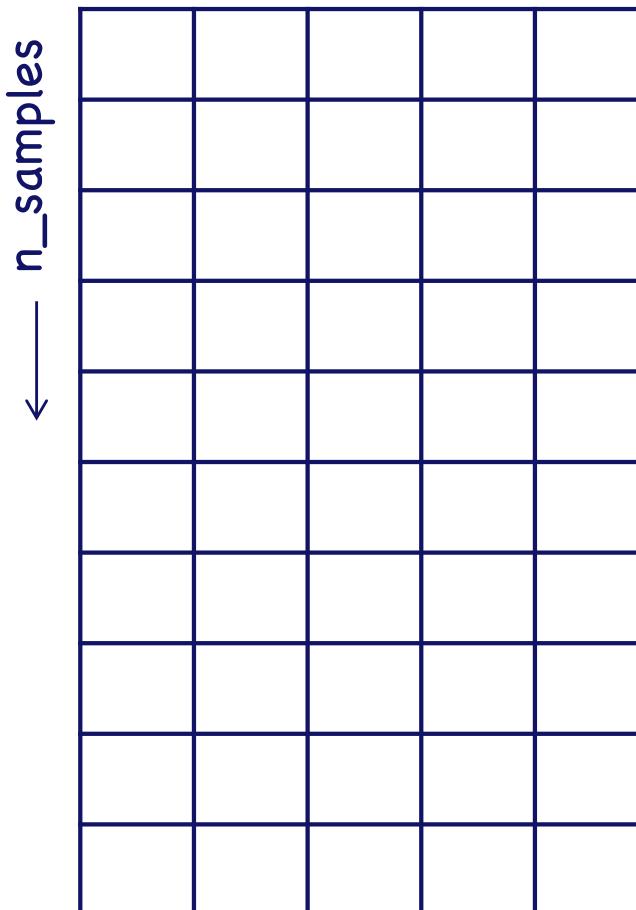
Representation of data

Raw data



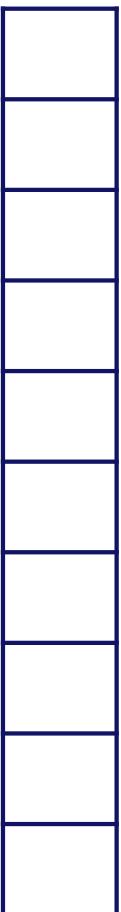
Feature matrix (X)

`n_features` →

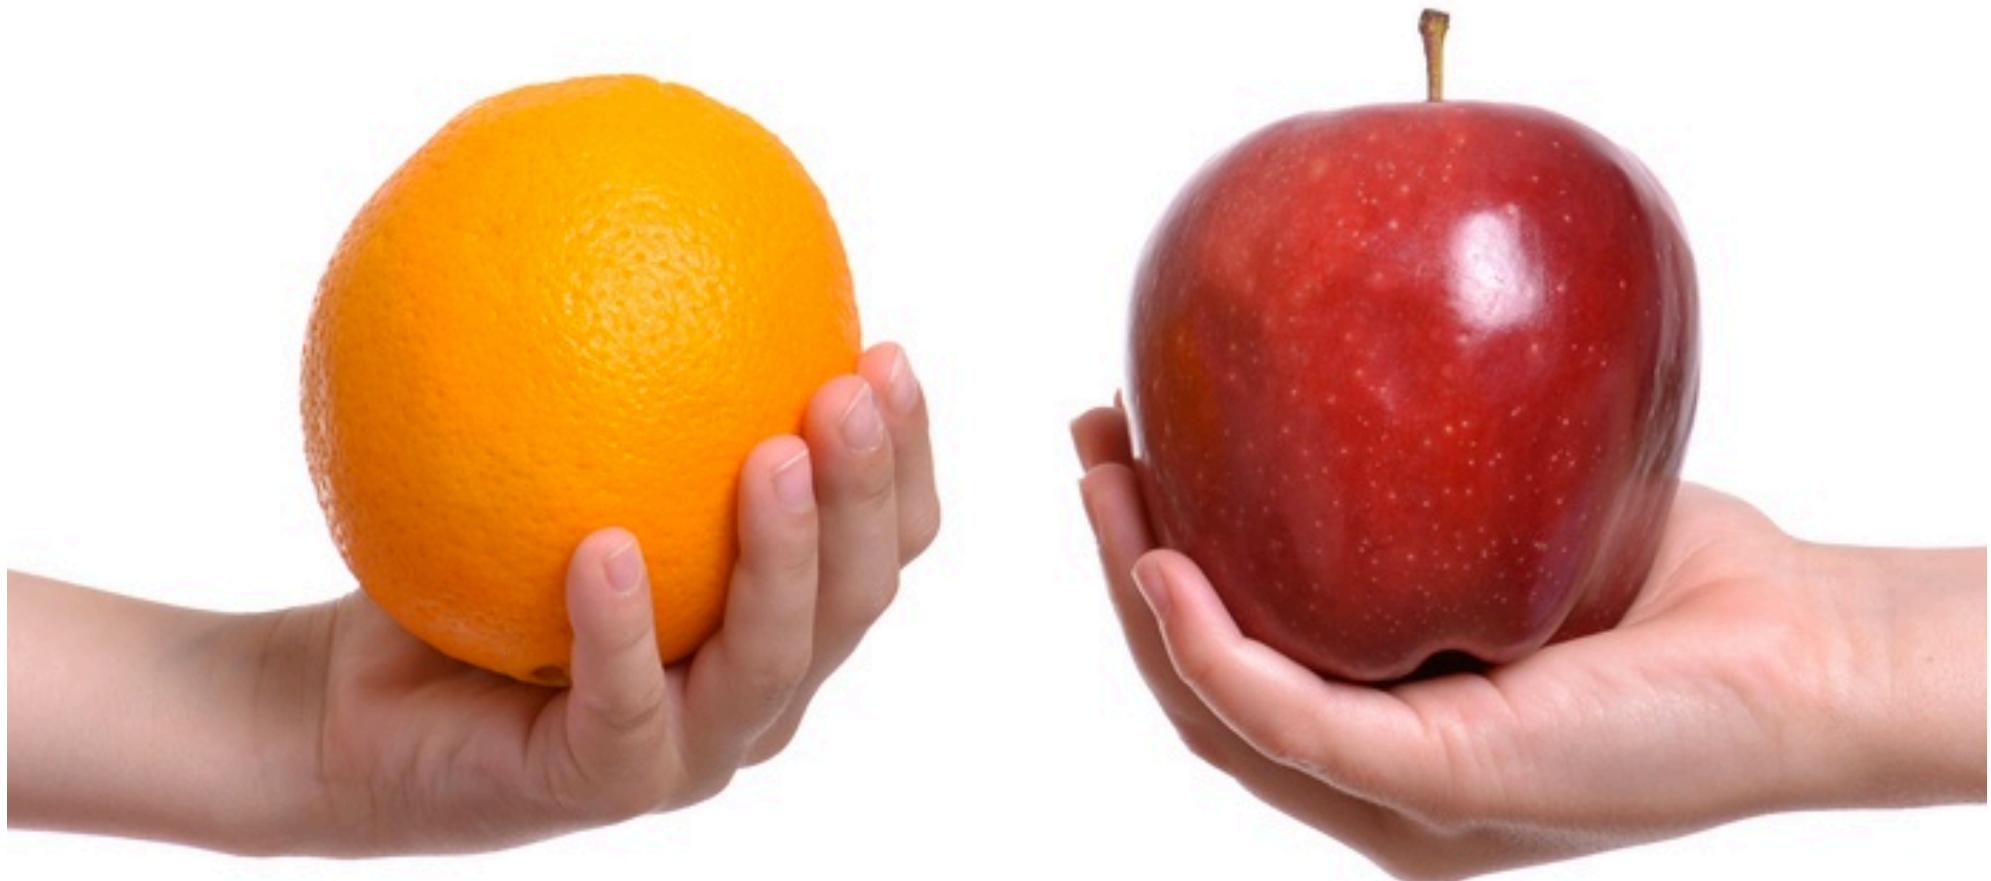


Target (y)

n_samples



Example



Example



Feature matrix (X)

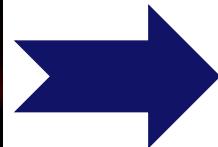
`n_features` →

A blank 8x8 grid for drawing or writing.

Target (y)

`n_samples`

Example



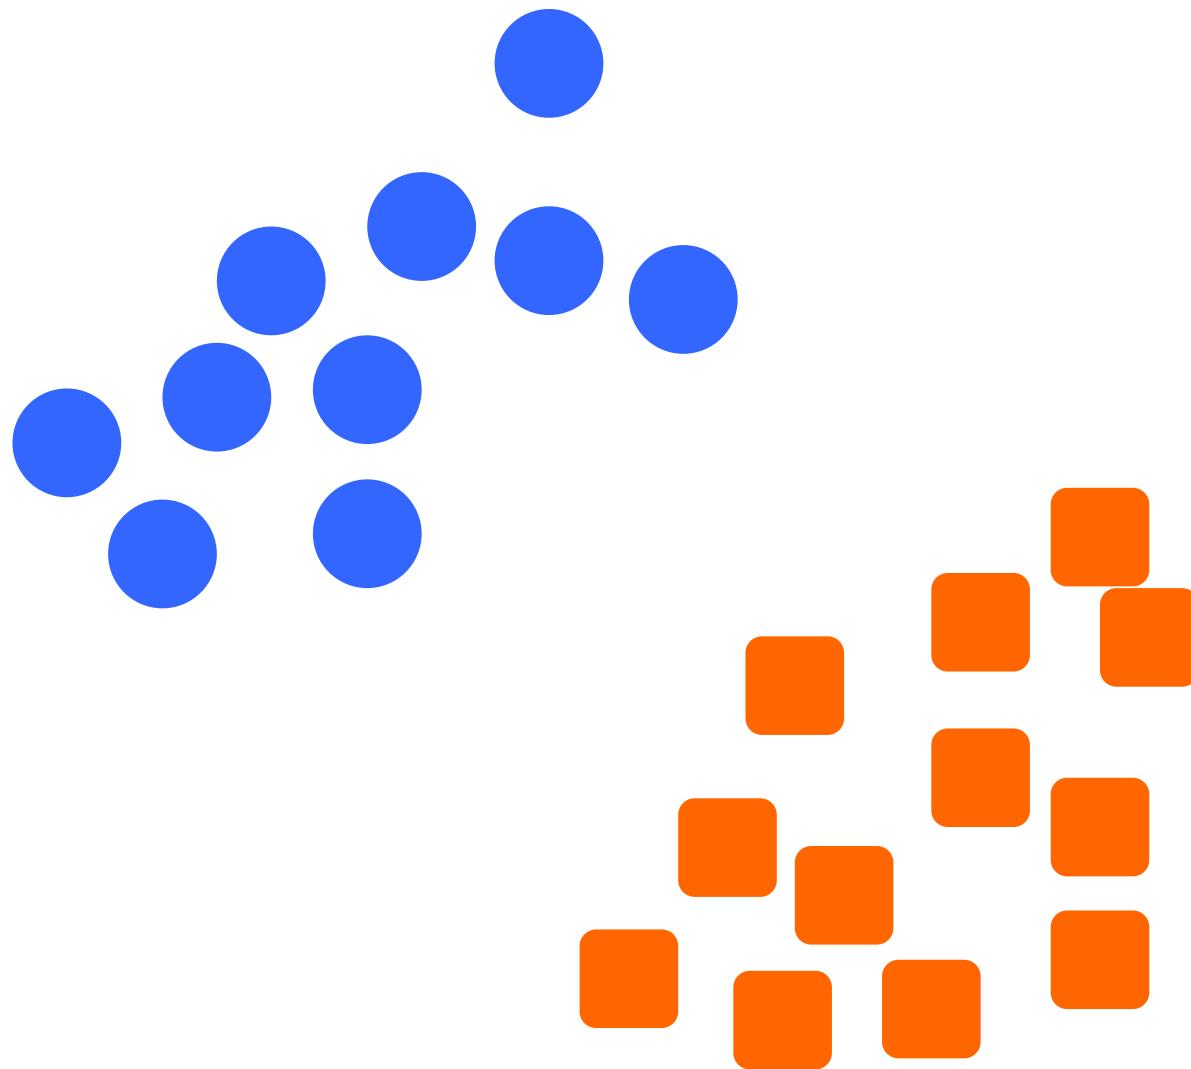
Feature matrix (X)

n_features →

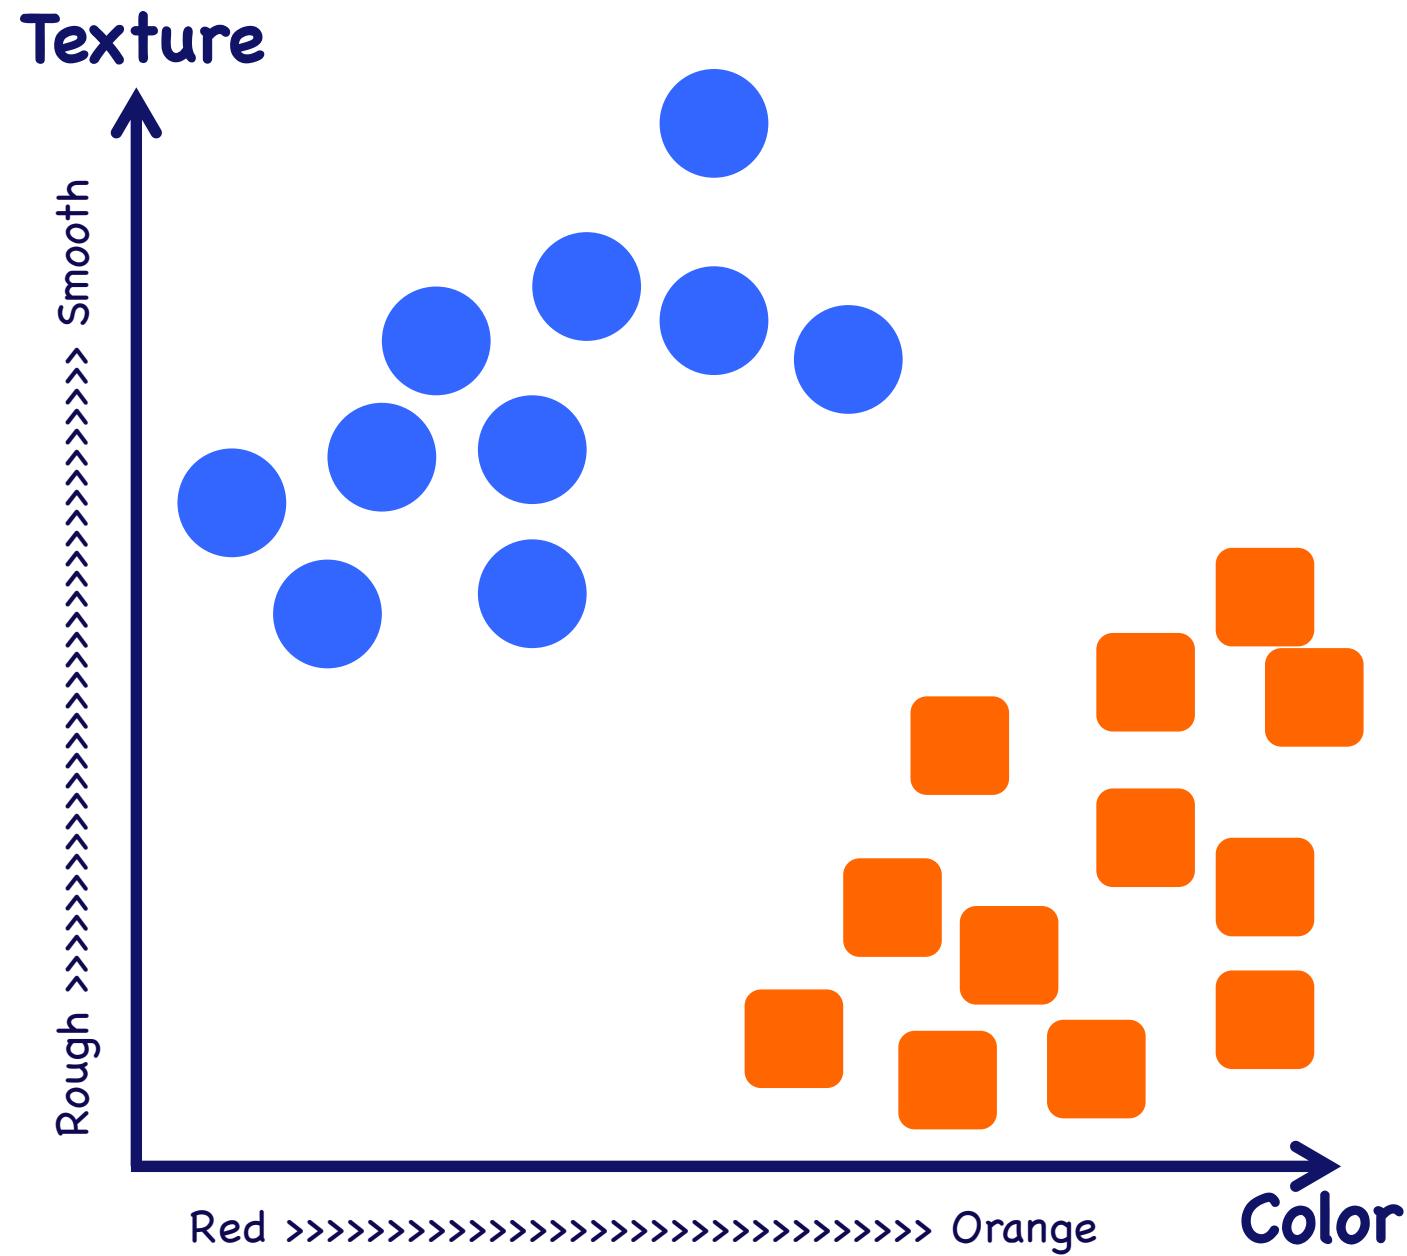
Target (y)

0
1
1
0
1
1
1
0
0
1

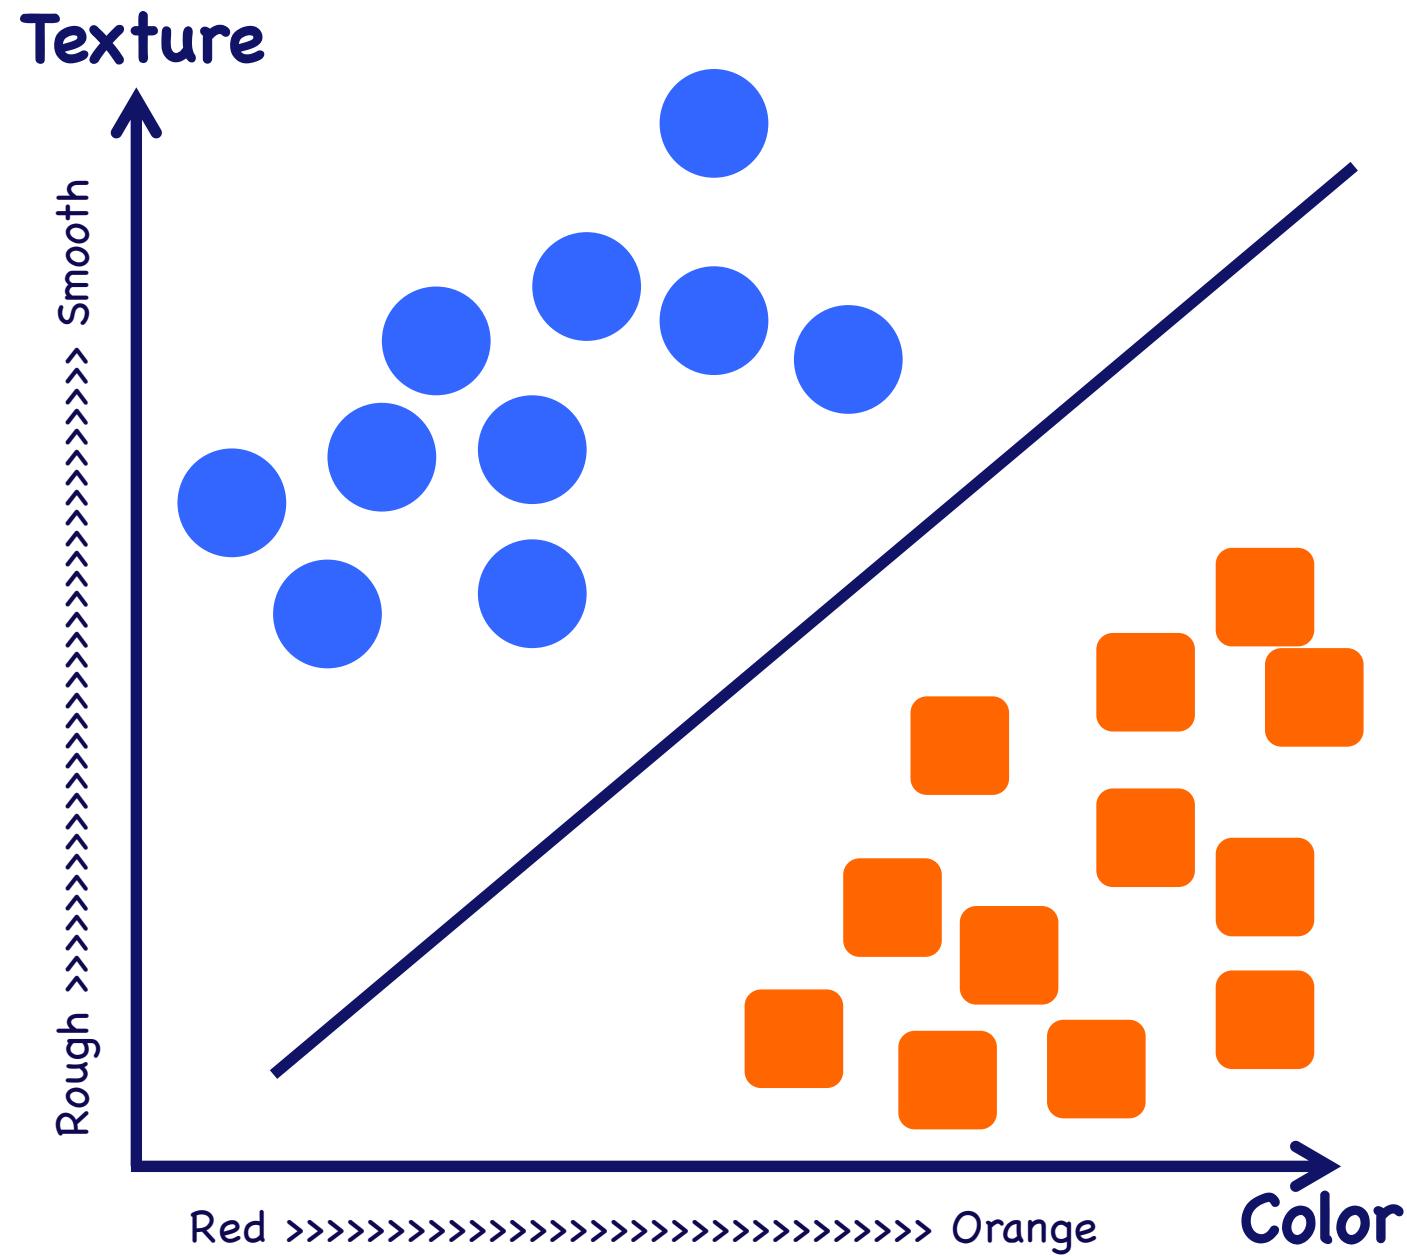
Classification



Classification



Classification

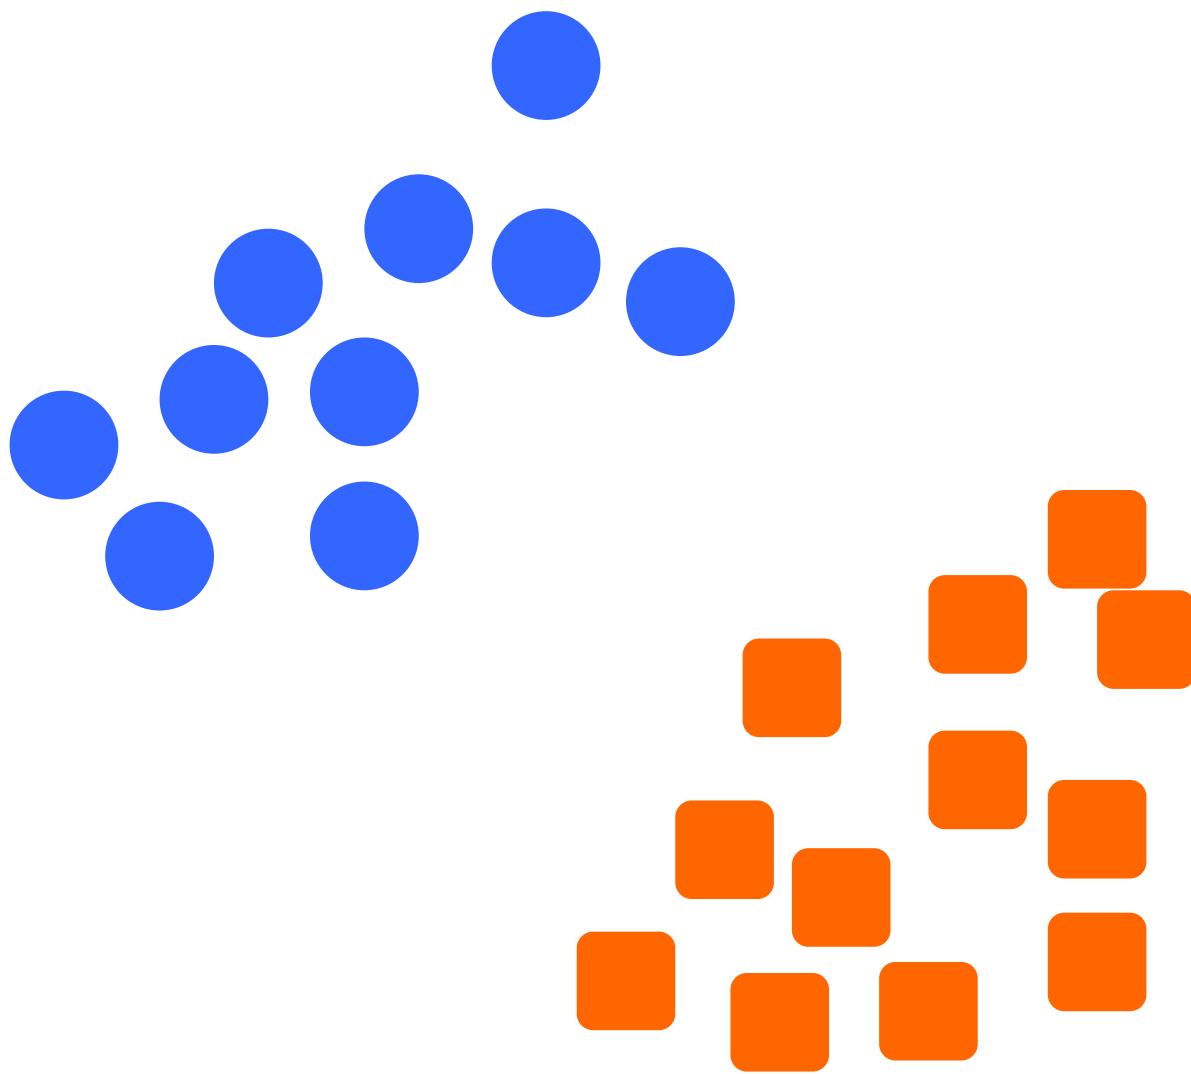


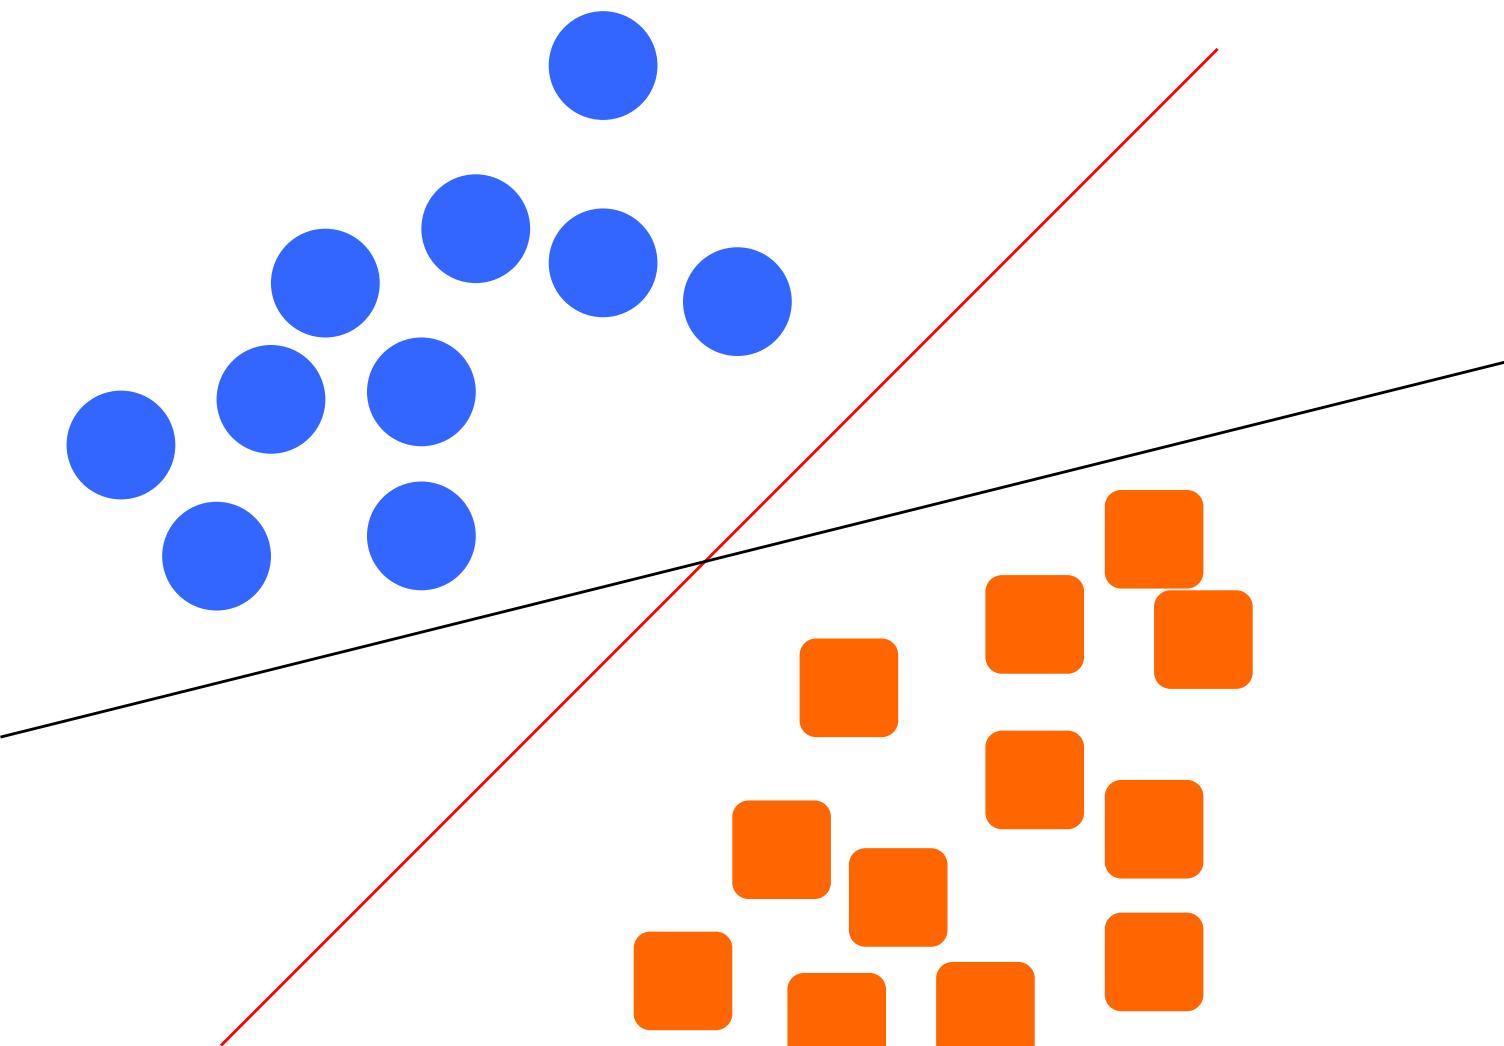
Classification

Support Vector Machine
Artificial Neural Network

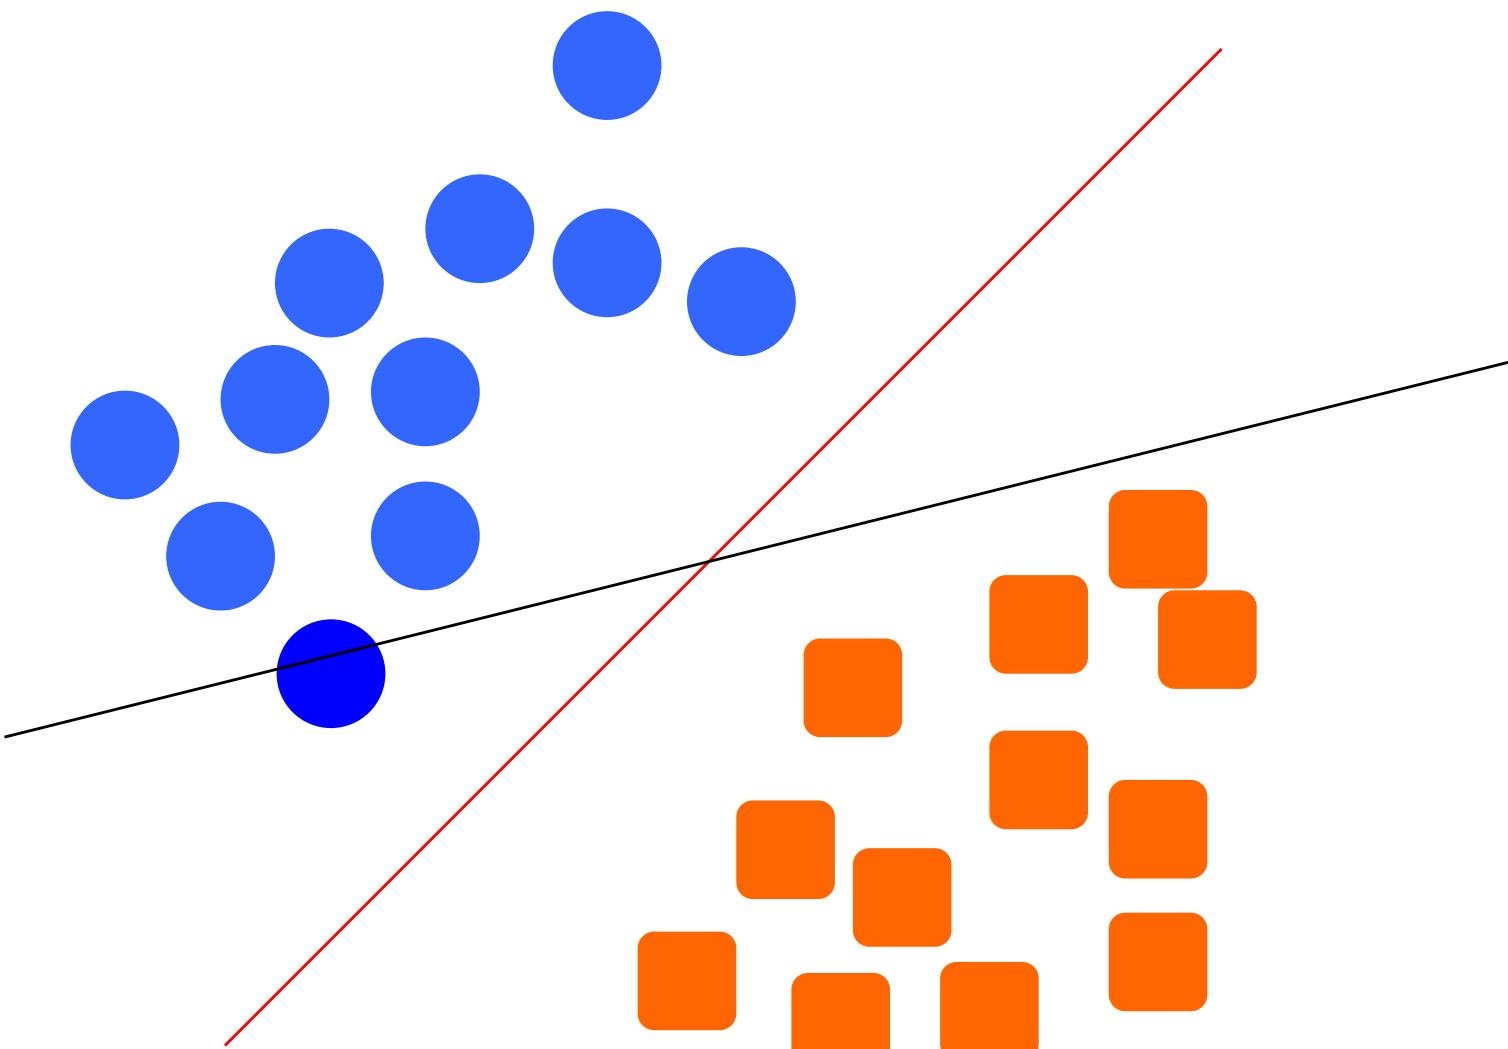


Supervised
learning

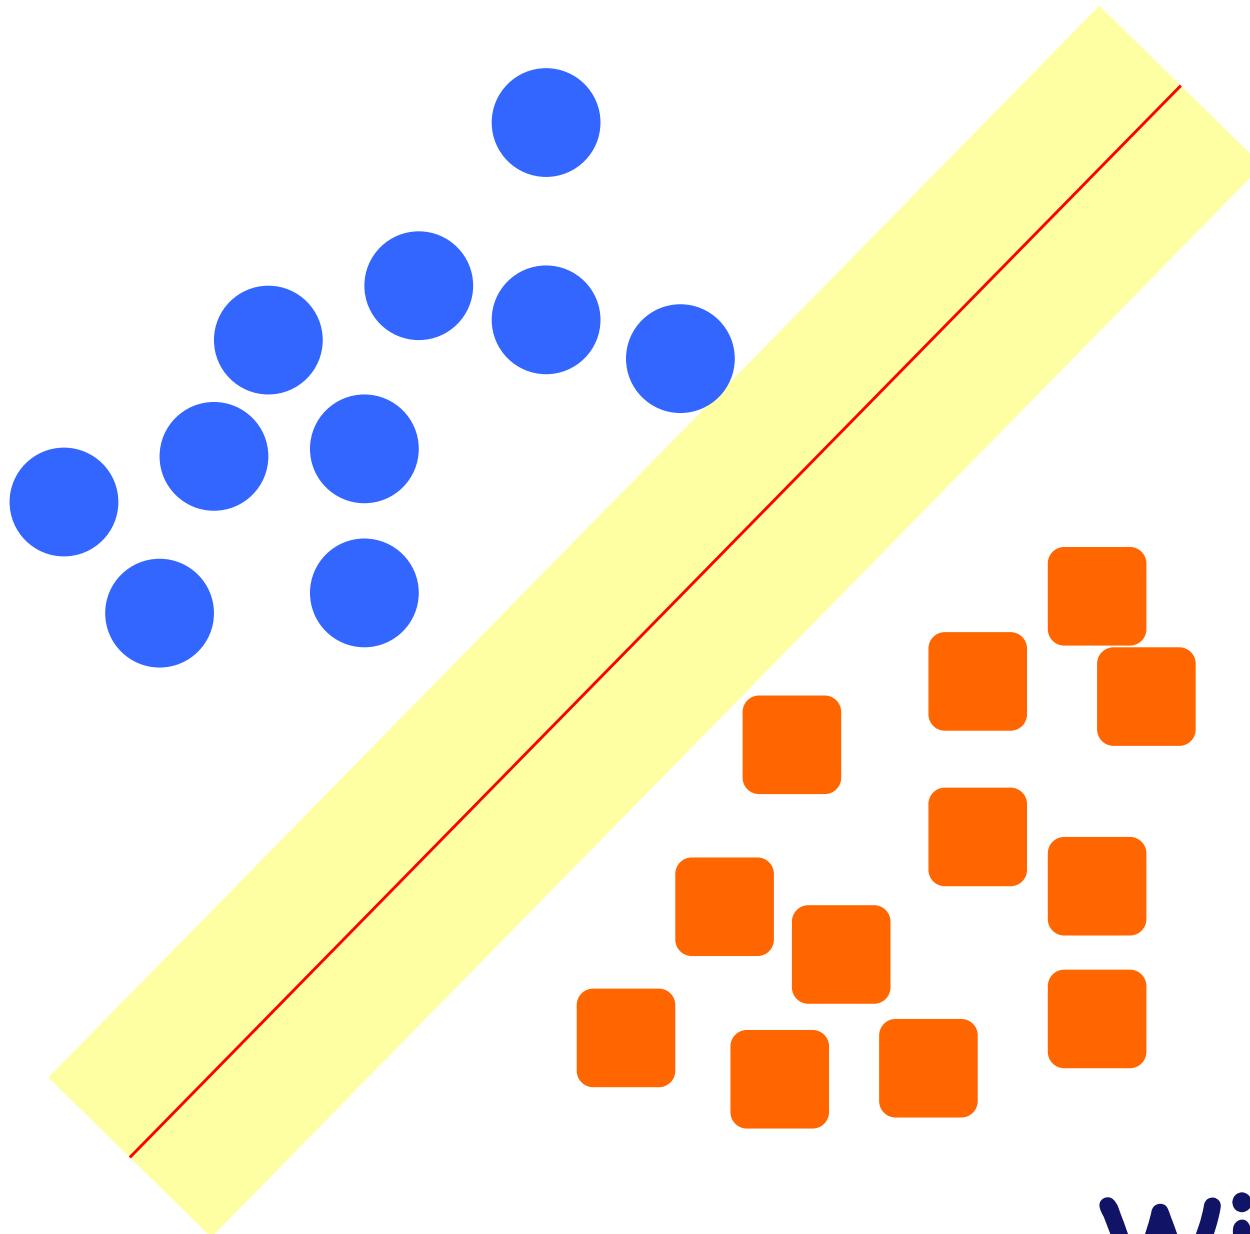




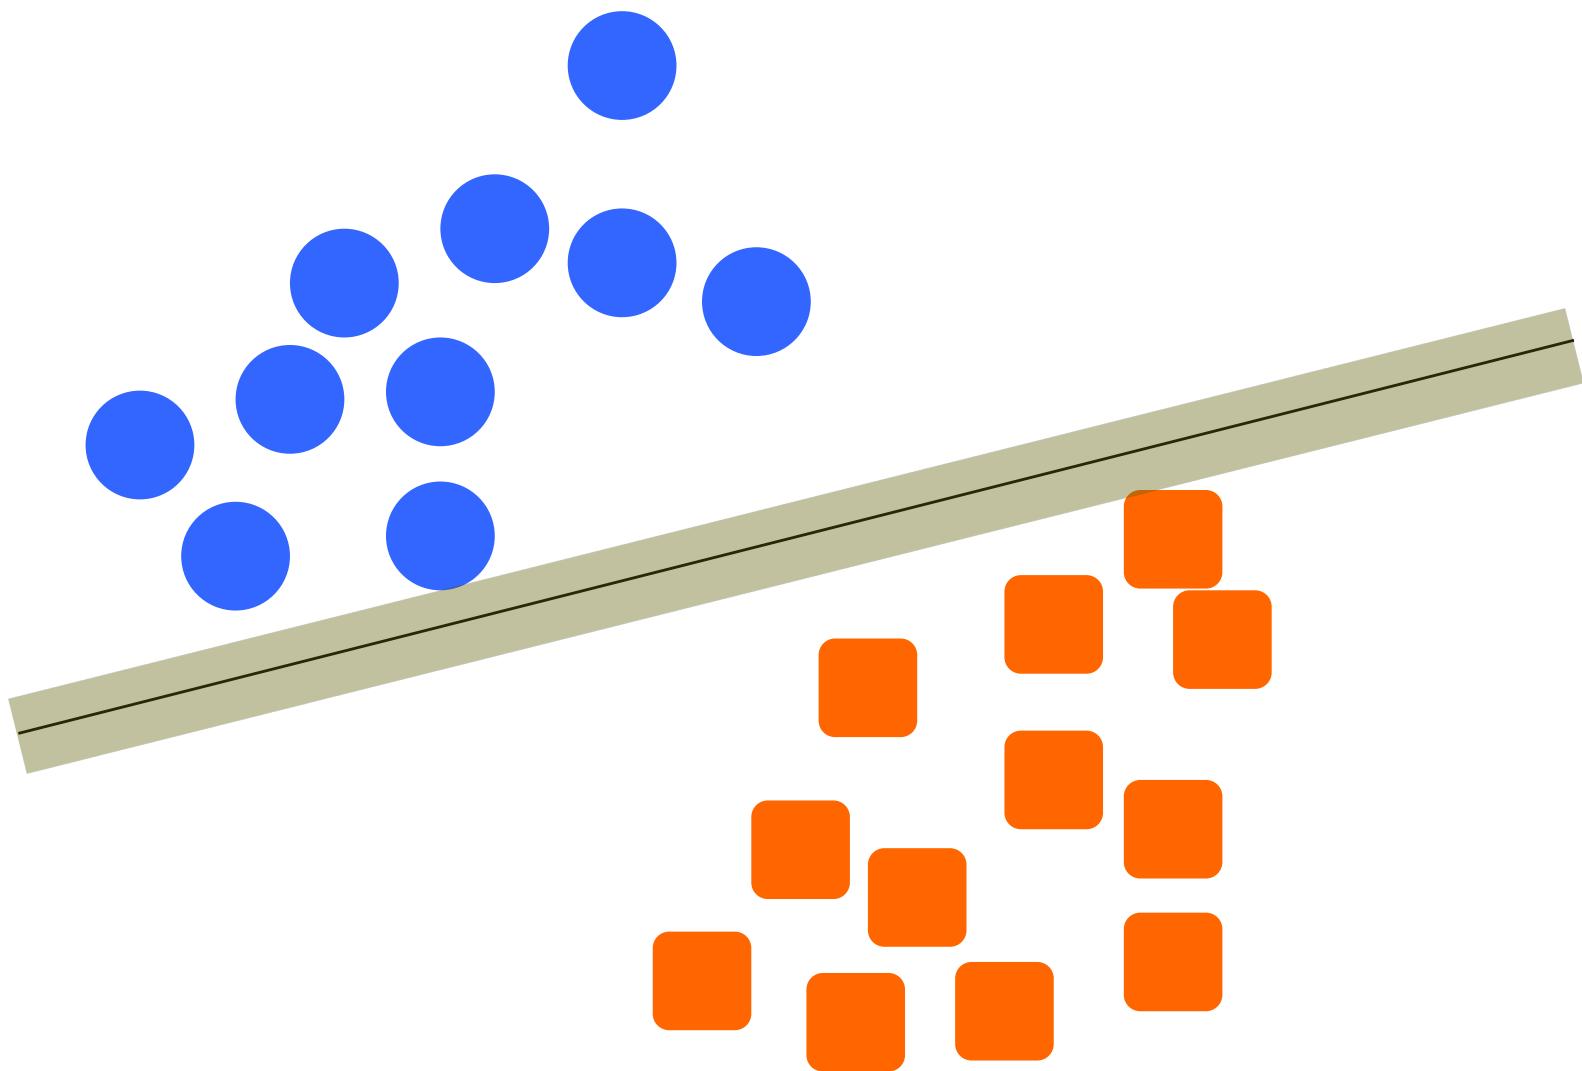
Which one is
better?



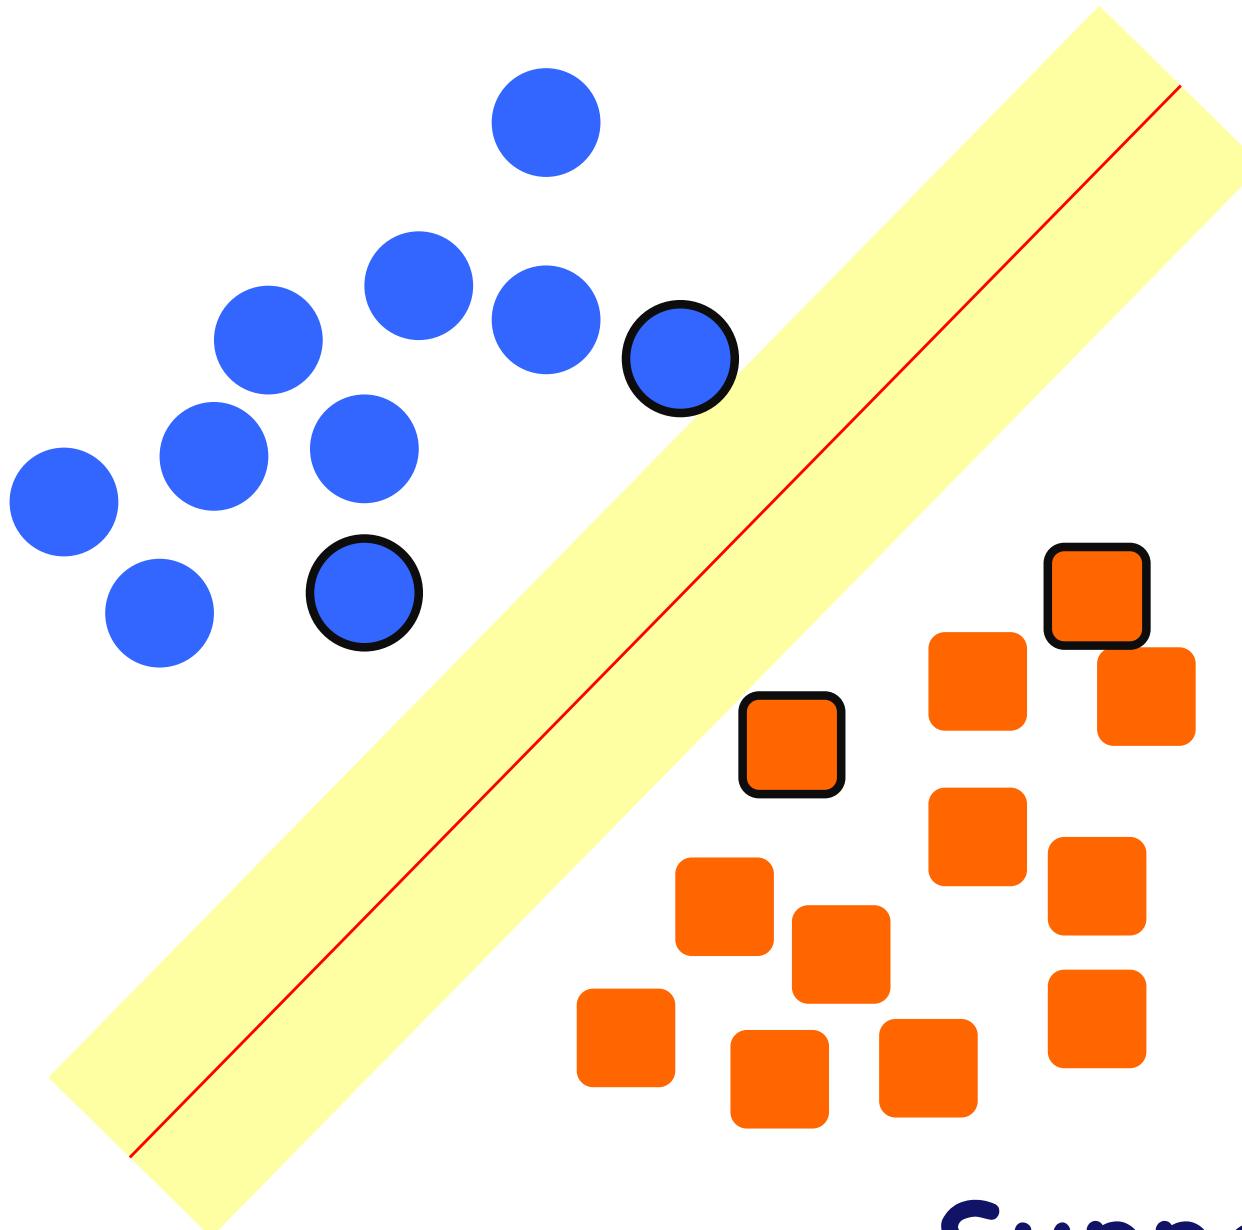
A new data point



Wide margin



Narrow margin

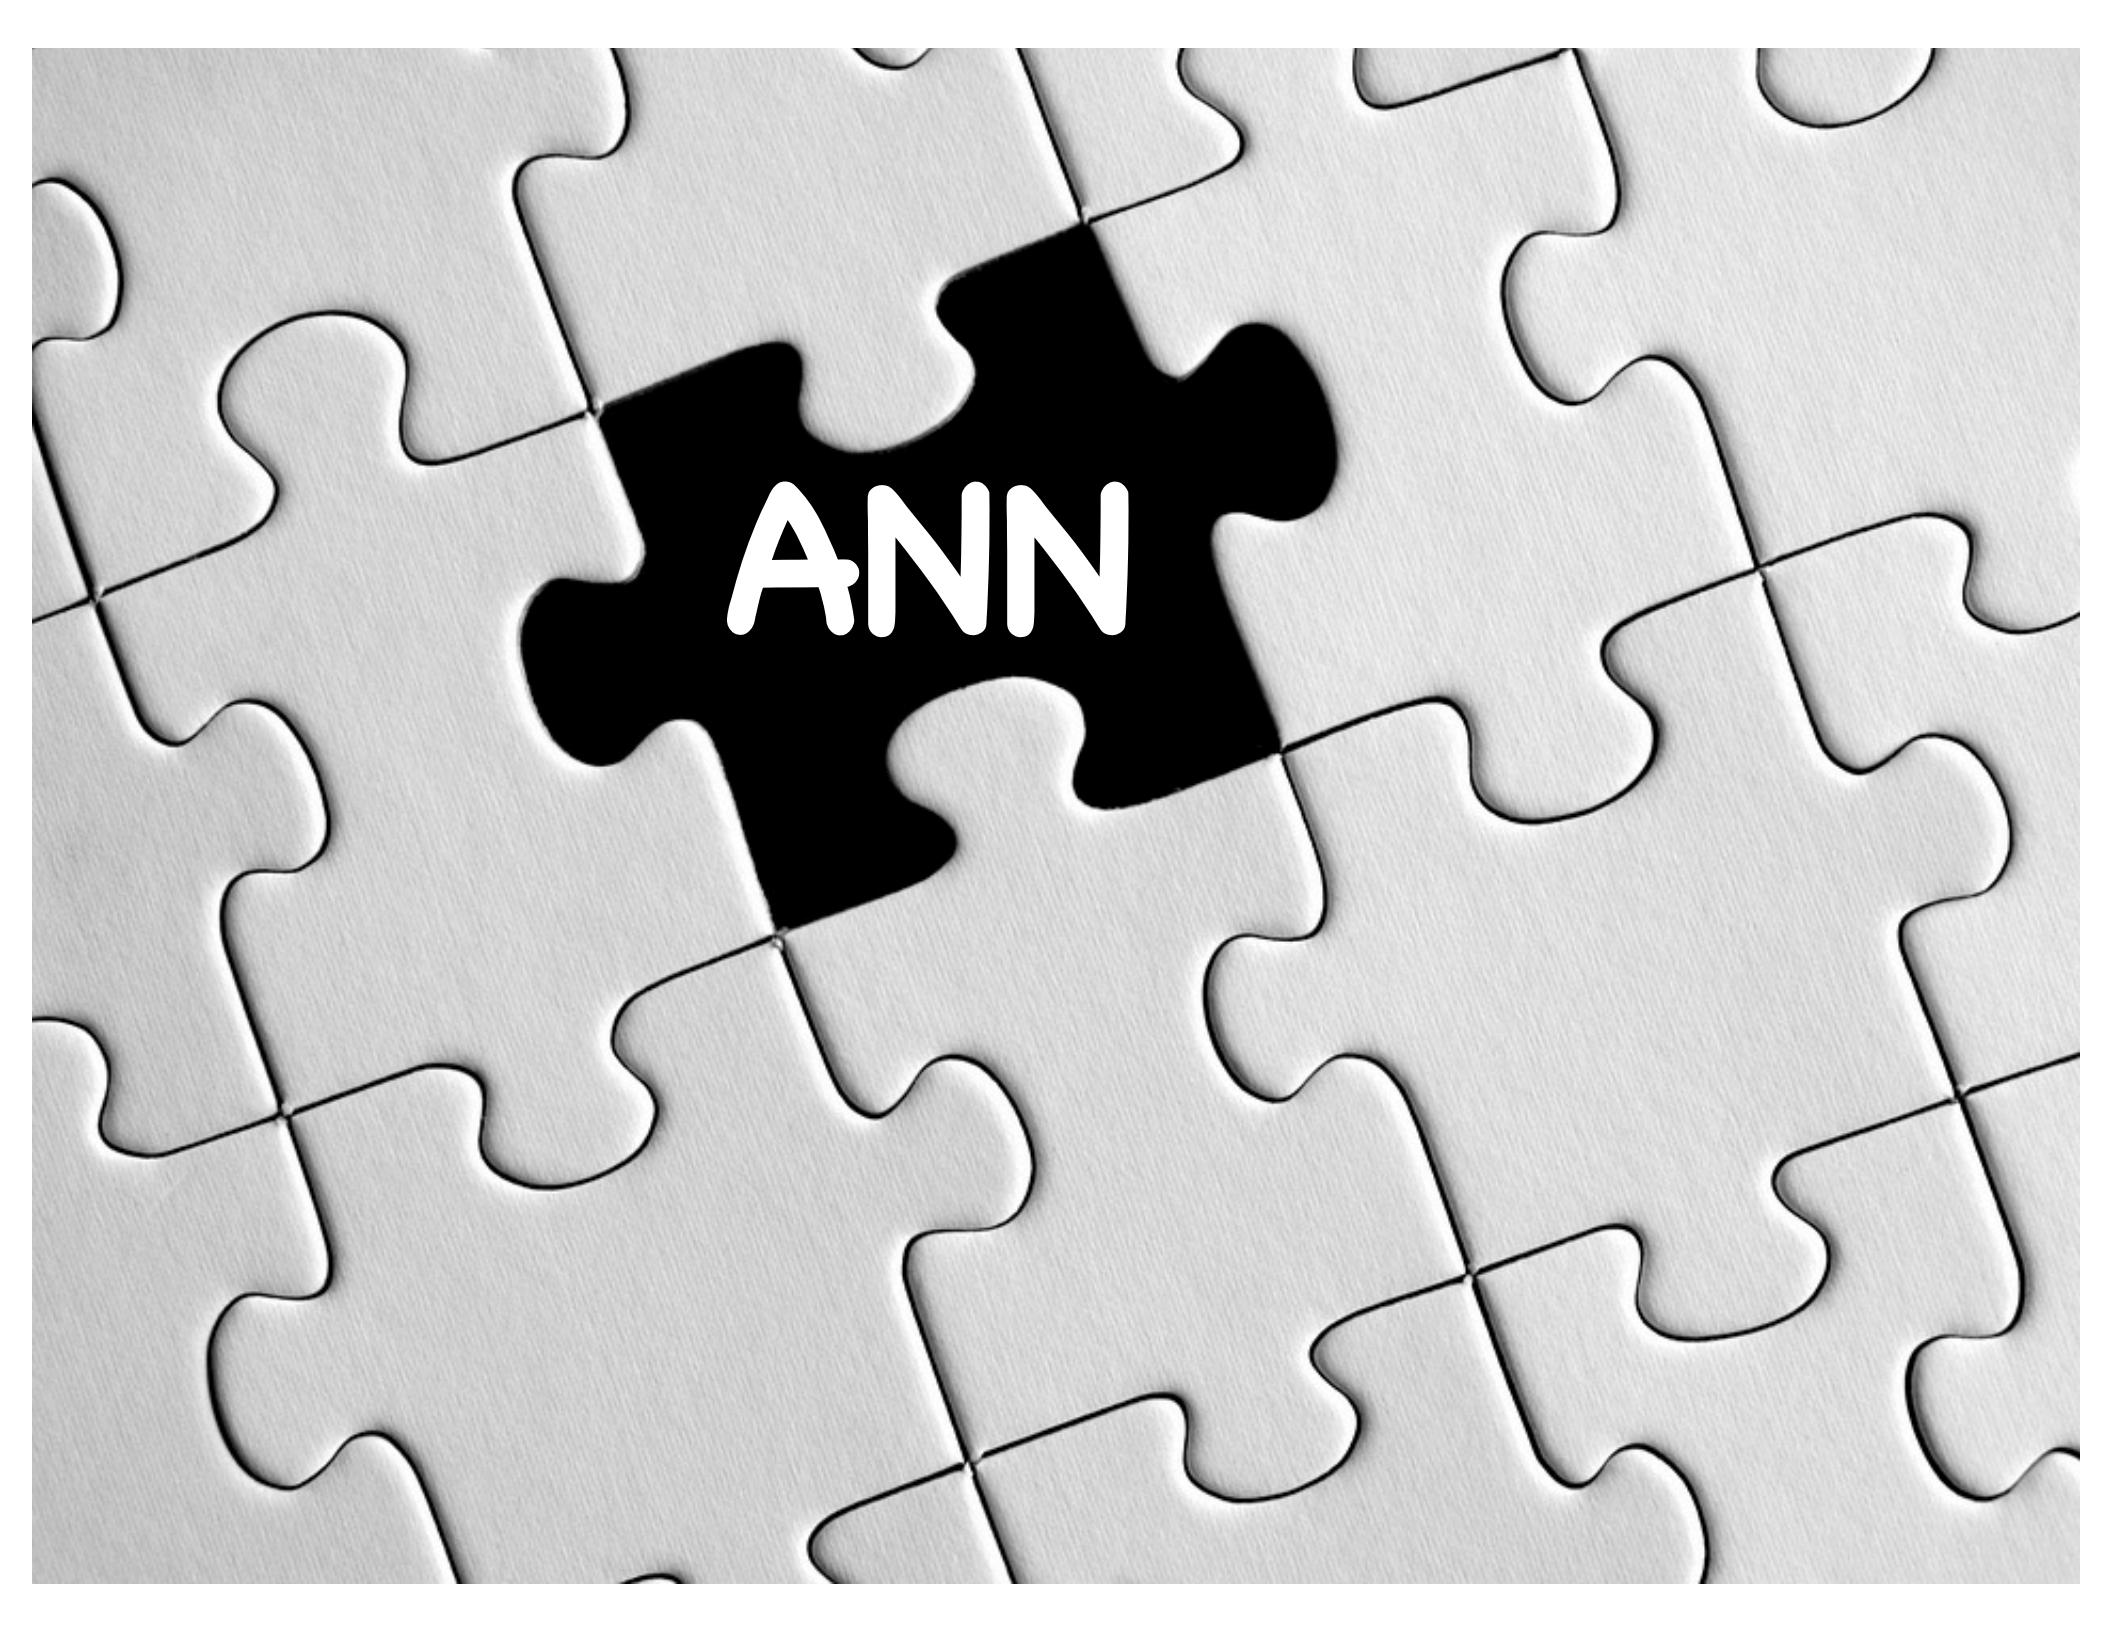


Support Vectors

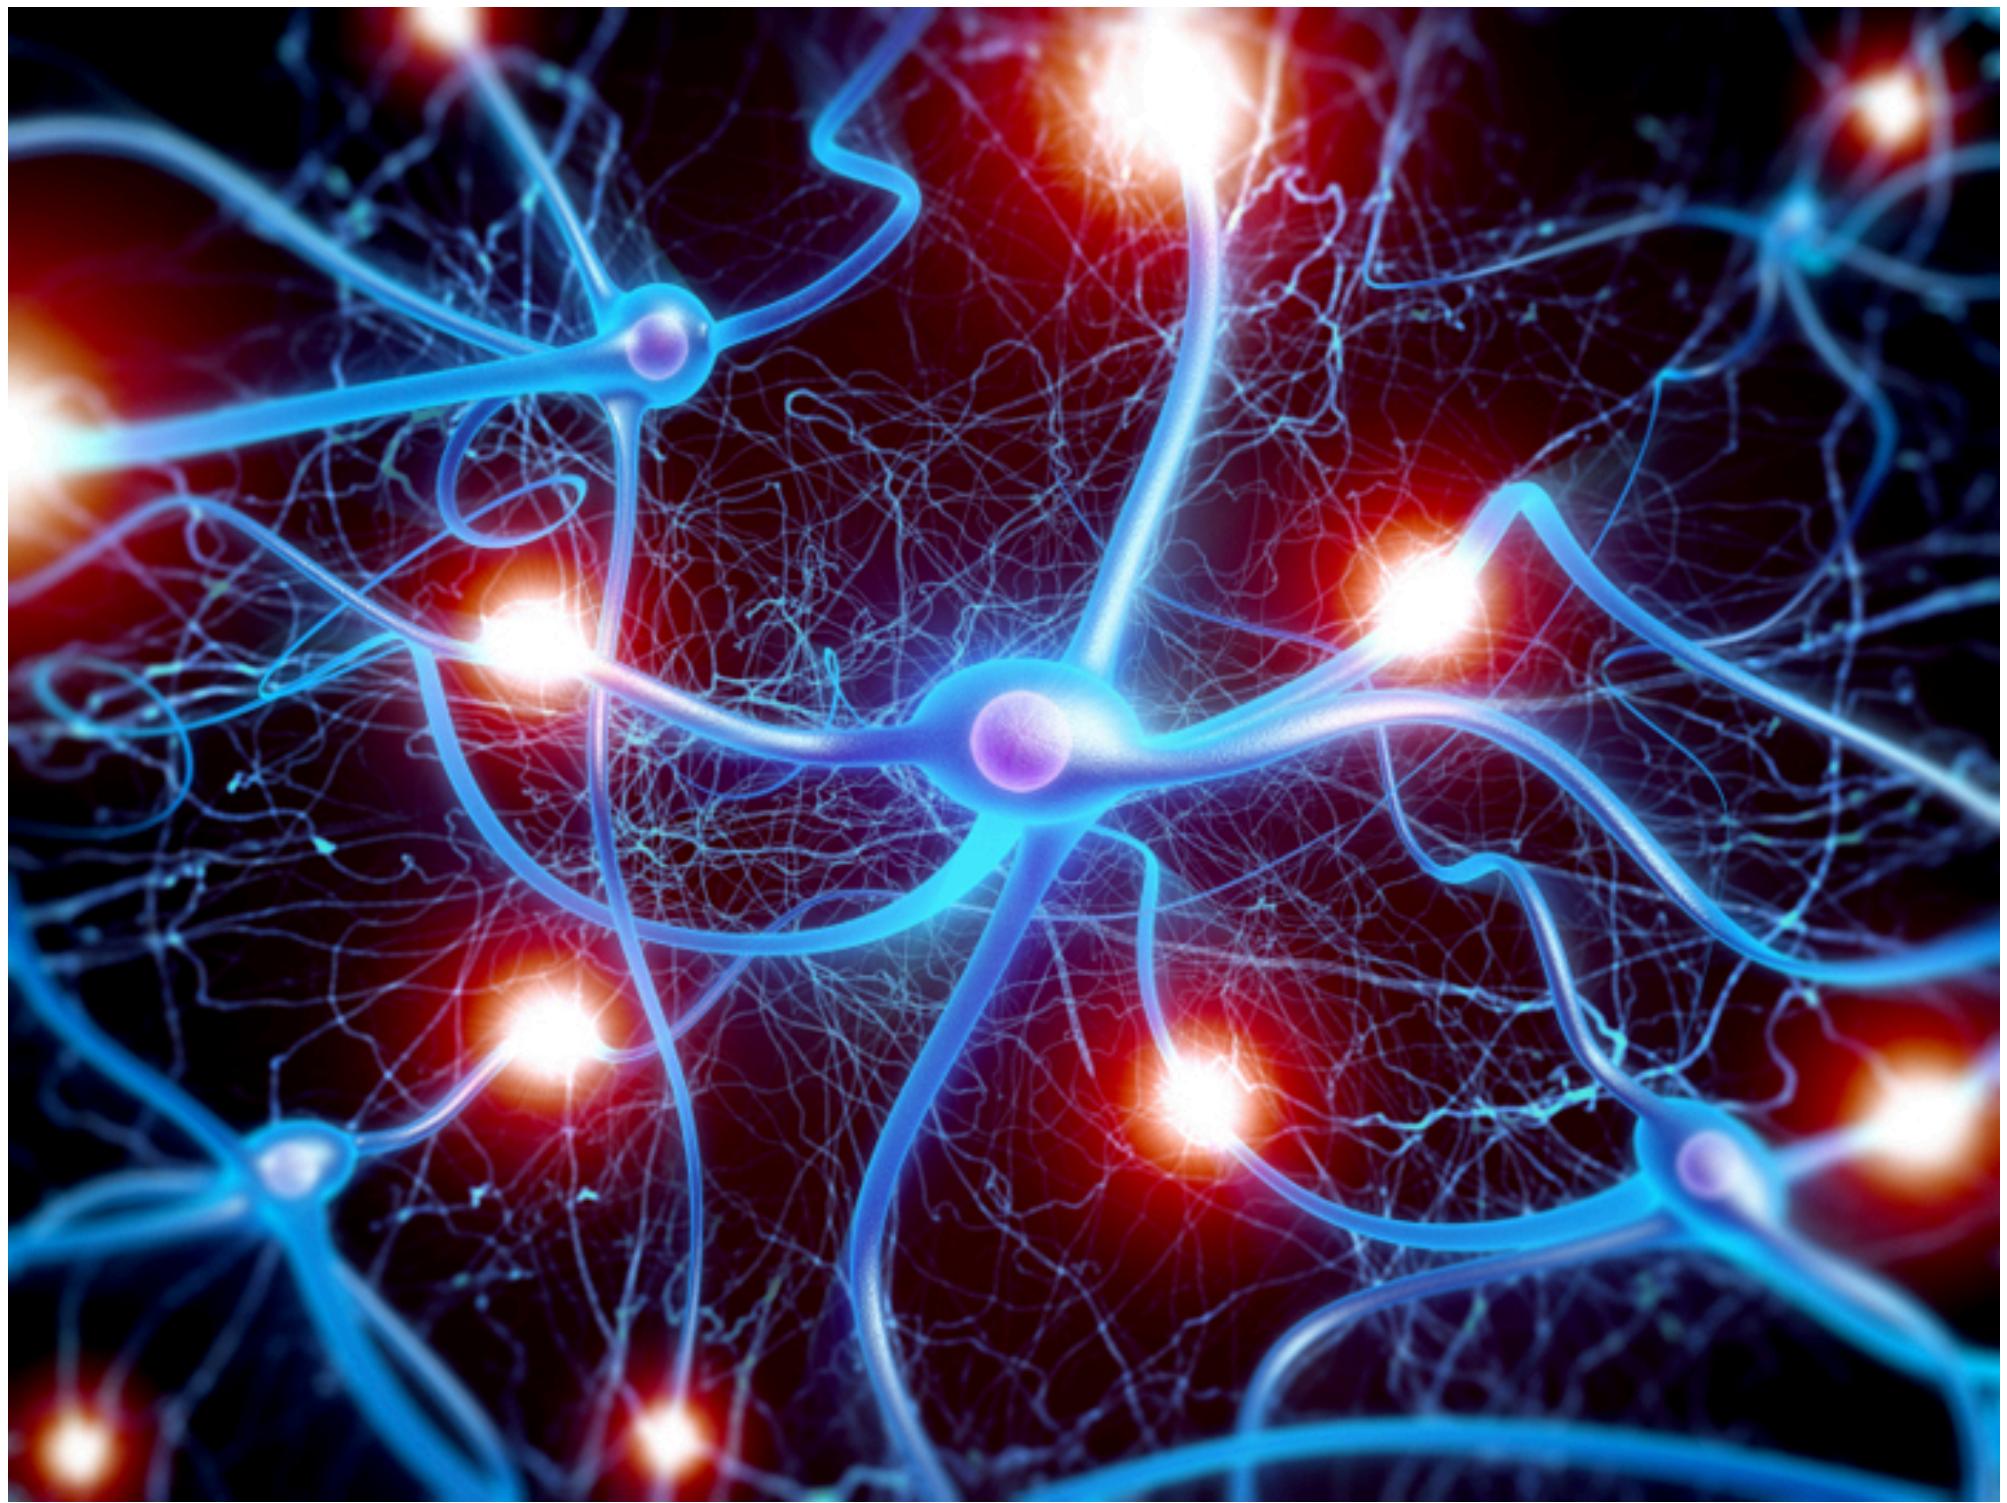
It is maximize the margin!

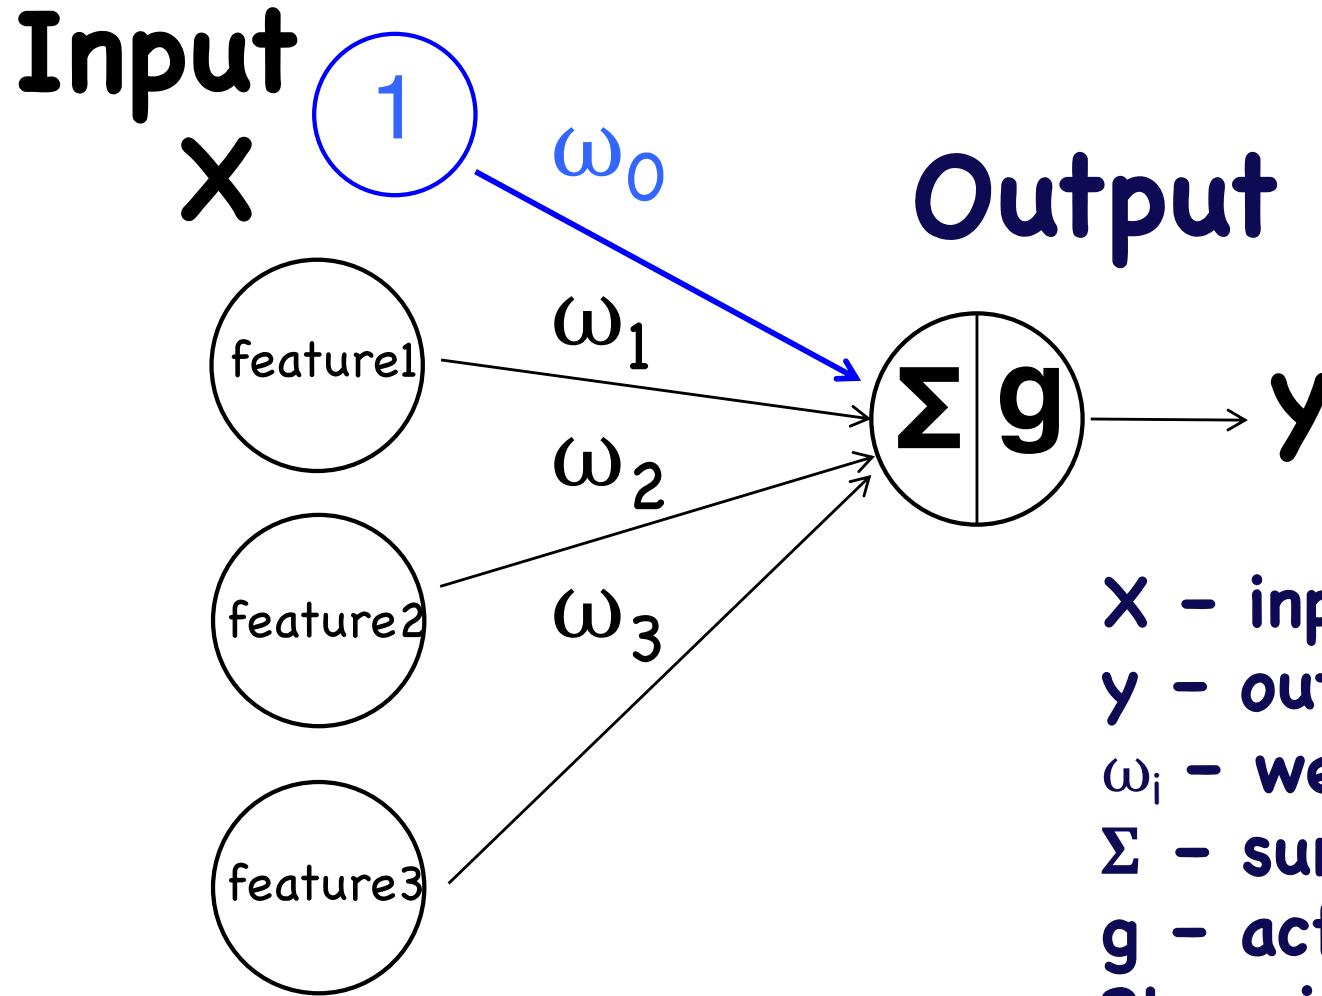
Support Vectors

More details: <https://www.youtube.com/watch?v=eHsErIPJWUU>



ANN



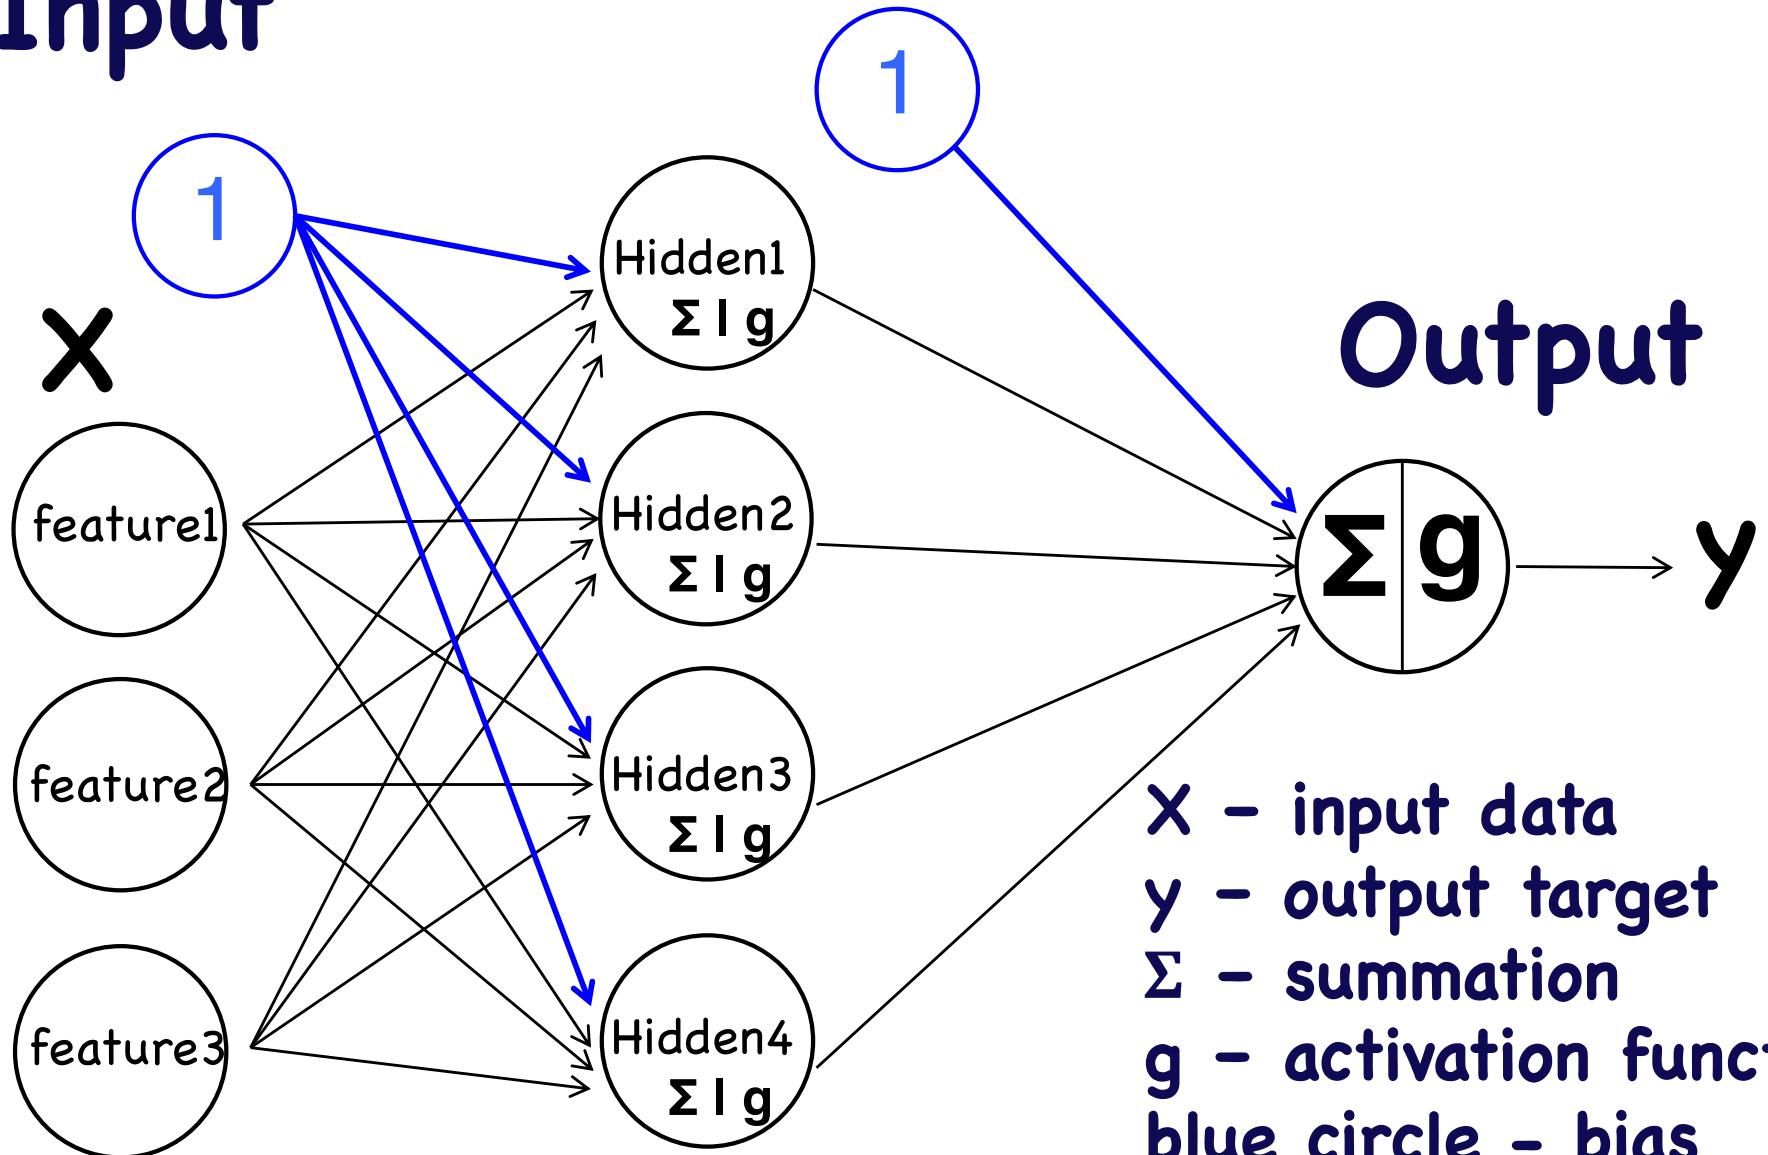


x - input data
 y - output target
 ω_i - weights
 Σ - summation
 g - activation function
 Blue circle - bias

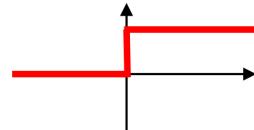
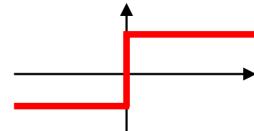
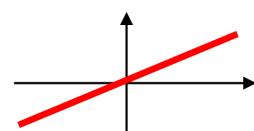
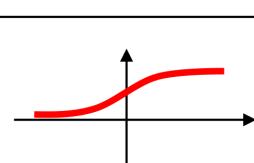
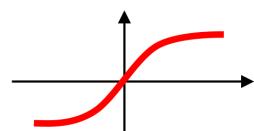
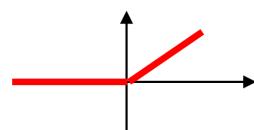
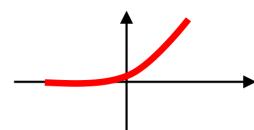
$$Z = \Sigma = \omega_0 x_0 + \omega_1 x_1 + \omega_2 x_2 + \omega_3 x_3 + \dots + \omega_n x_n$$

$$y = g(\omega_0 x_0 + \omega_1 x_1 + \omega_2 x_2 + \omega_3 x_3 + \dots + \omega_n x_n)$$

Input



X - input data
y - output target
 Σ - summation
g - activation function
blue circle - bias

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	



YOU'RE IN MY SPOT

GRAPHICS GARAGE

Input



.

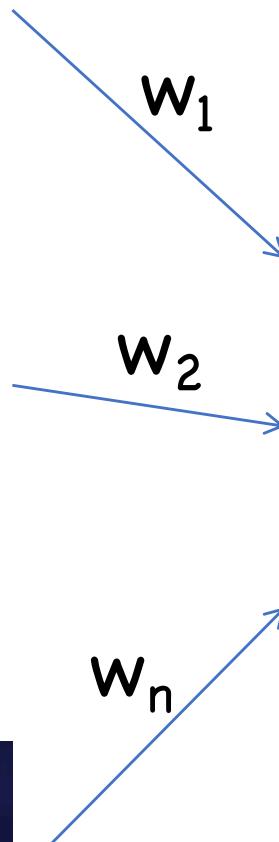
.

.

.



Intuitive Artificial Neural Network



$$F(\text{eye} \times w_1 + \text{nose} \times w_2 + \dots + \text{mouth} \times w_n)$$



Output



Input



.

.

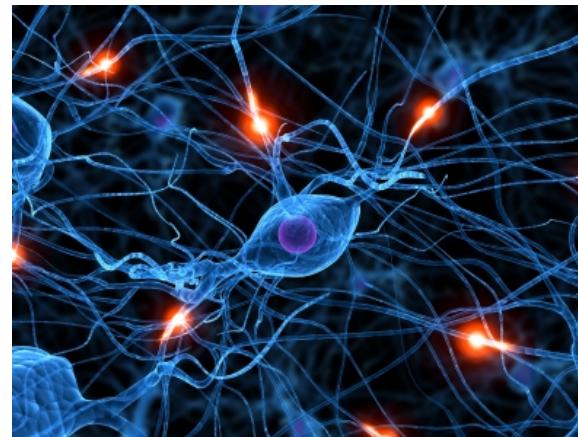
.



Intuitive Artificial Neural Network

Output

$$F(\text{eye} \times w_1 + \text{nose} \times w_2 + \dots + \text{mouth} \times w_n)$$



error
feedback



Input



.

.

.



Intuitive Artificial Neural Network

Output

$$F(\text{eye} \times w_1 + \text{nose} \times w_2 + \dots + \text{mouth} \times w_n)$$



error
feedback



more photos

Input



.

.

.



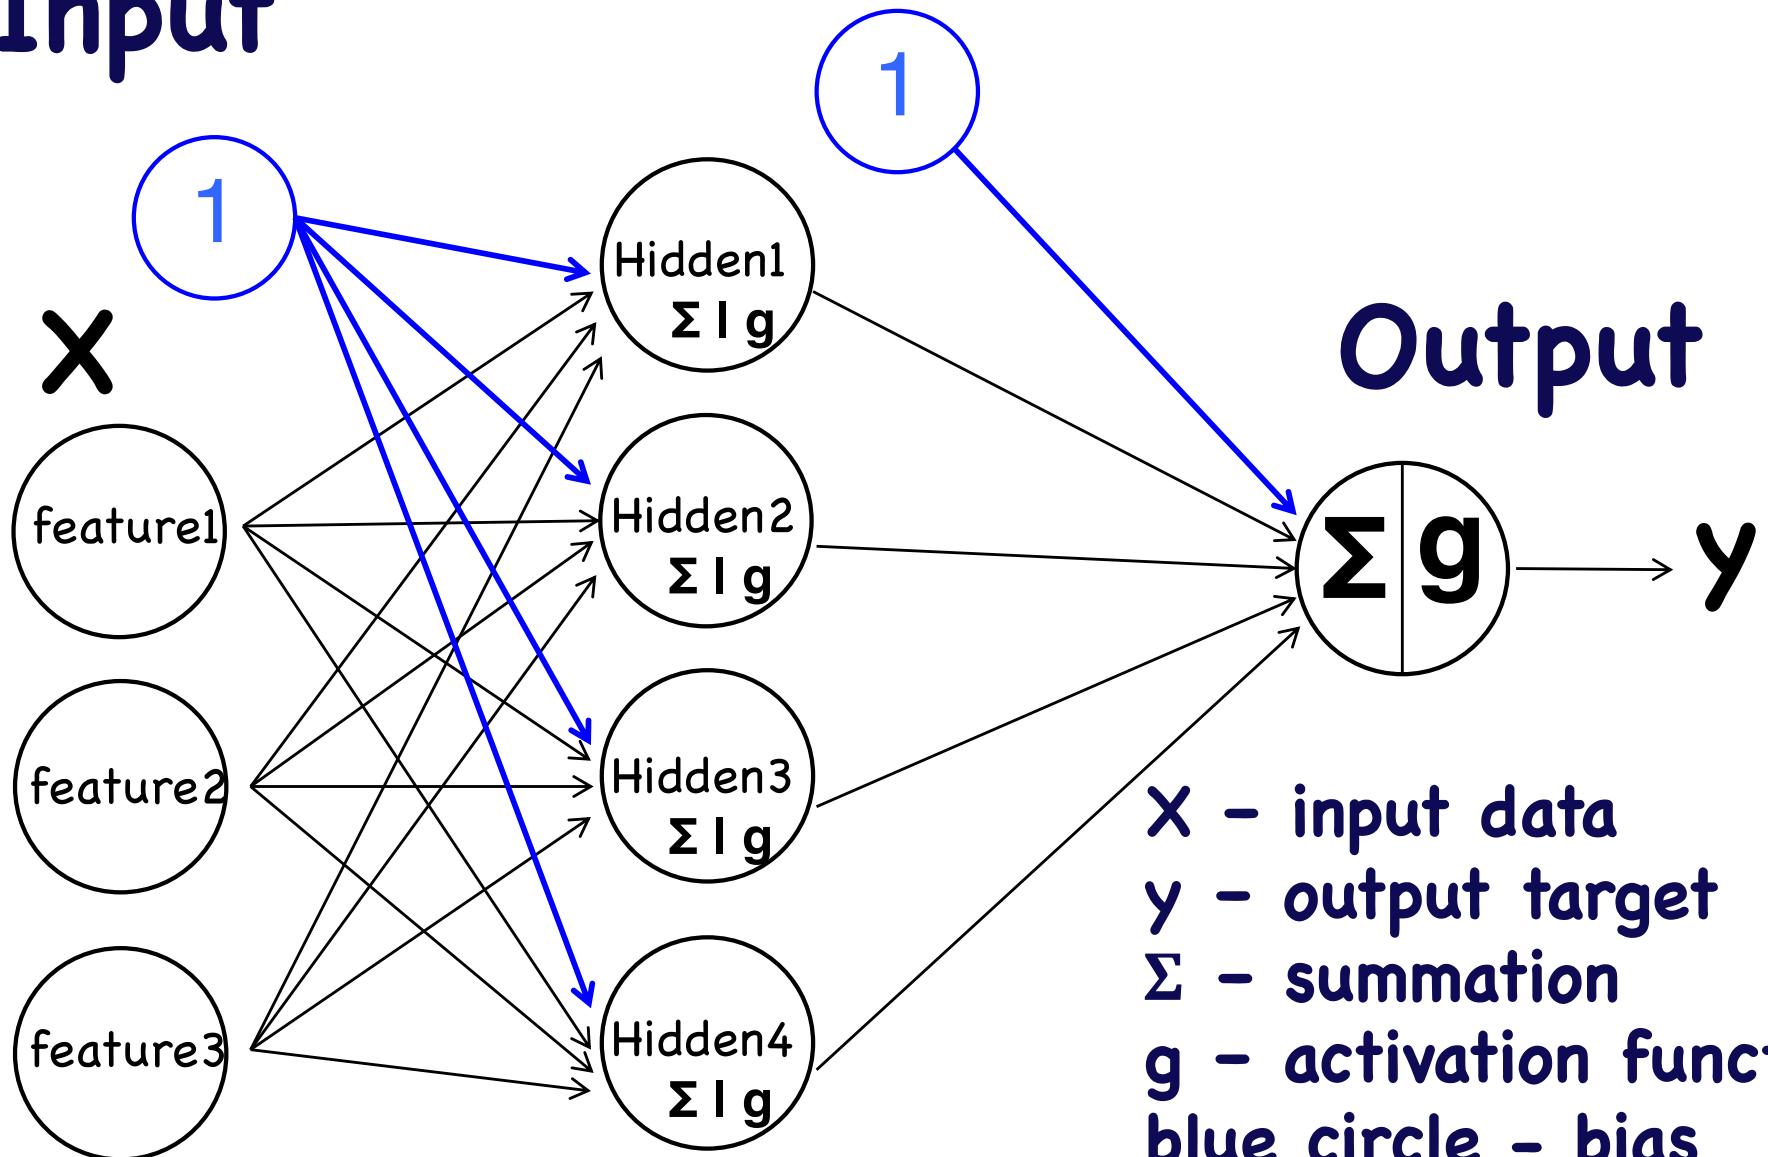
Intuitive Artificial Neural Network

Output

$$F(\text{eye} \times w_1 + \text{nose} \times w_2 + \dots + \text{mouth} \times w_n)$$



Input



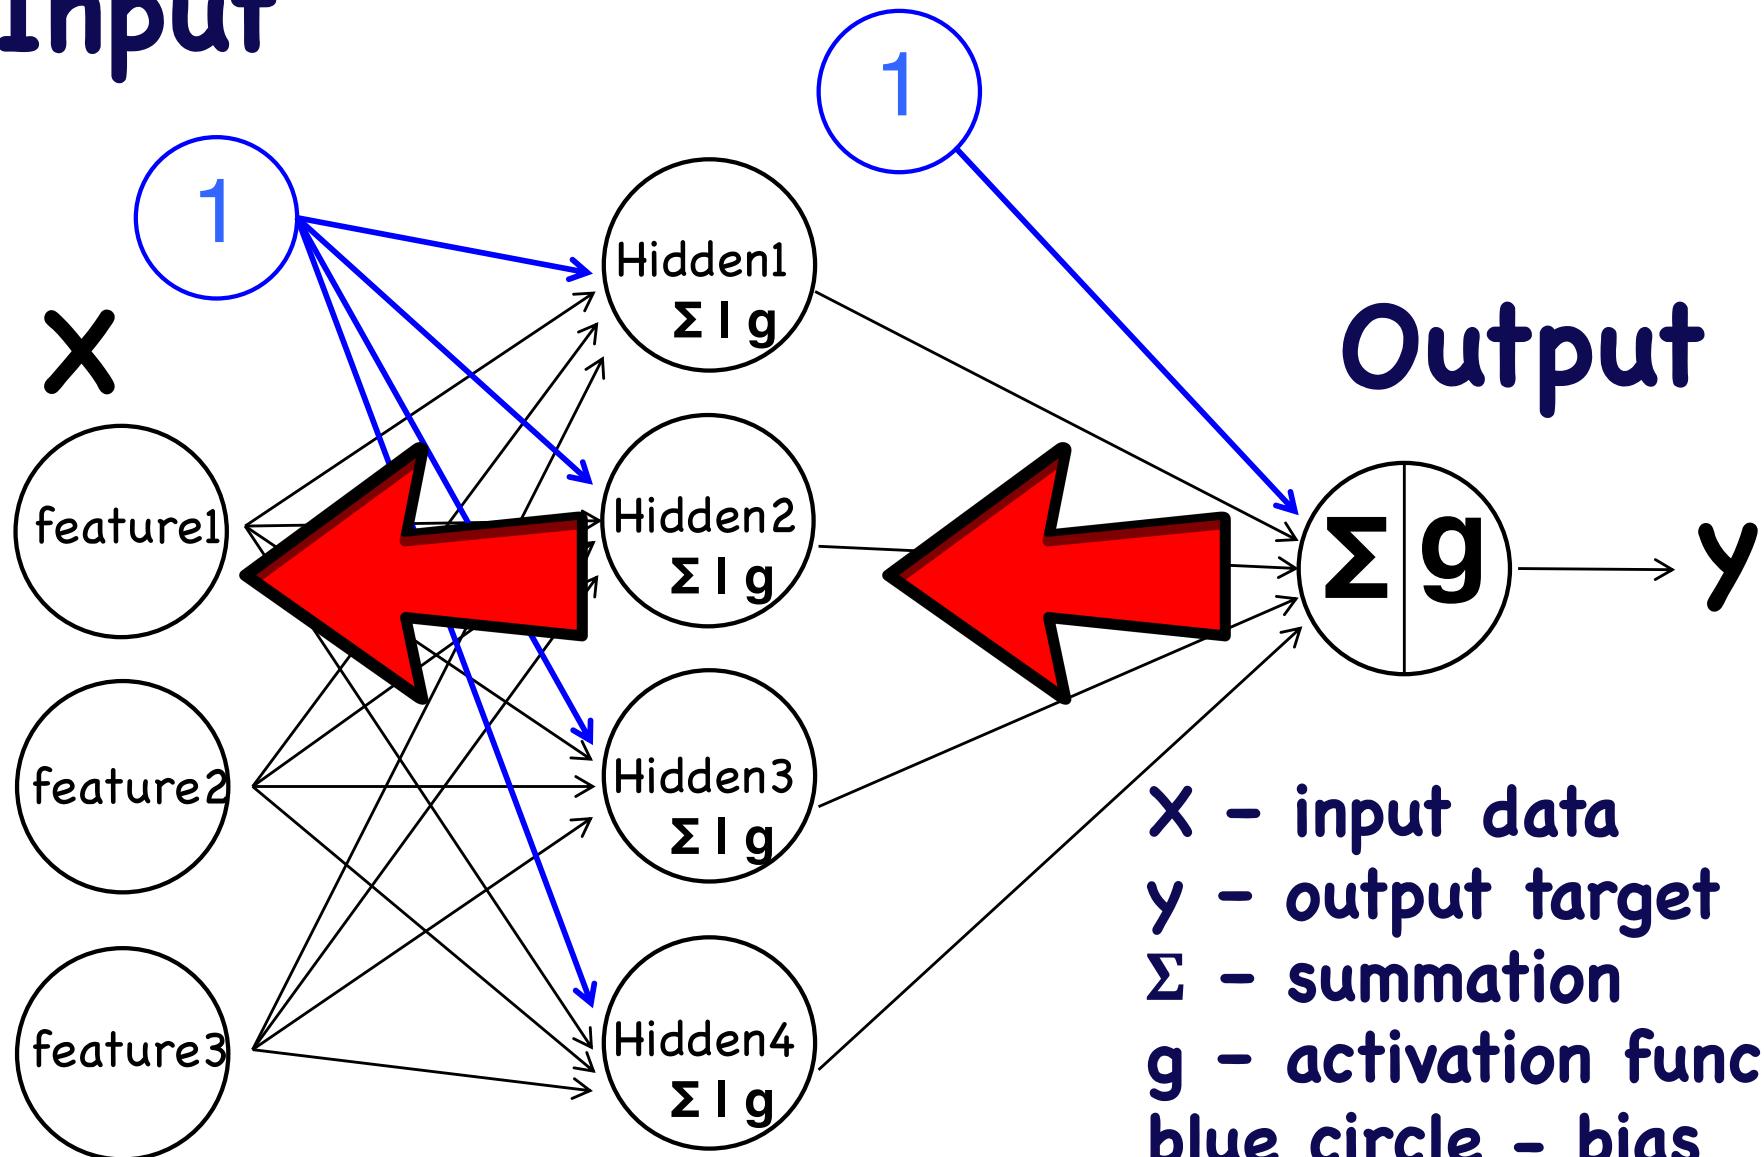
x - input data
y - output target
 Σ - summation
g - activation function
blue circle - bias

~~Error~~



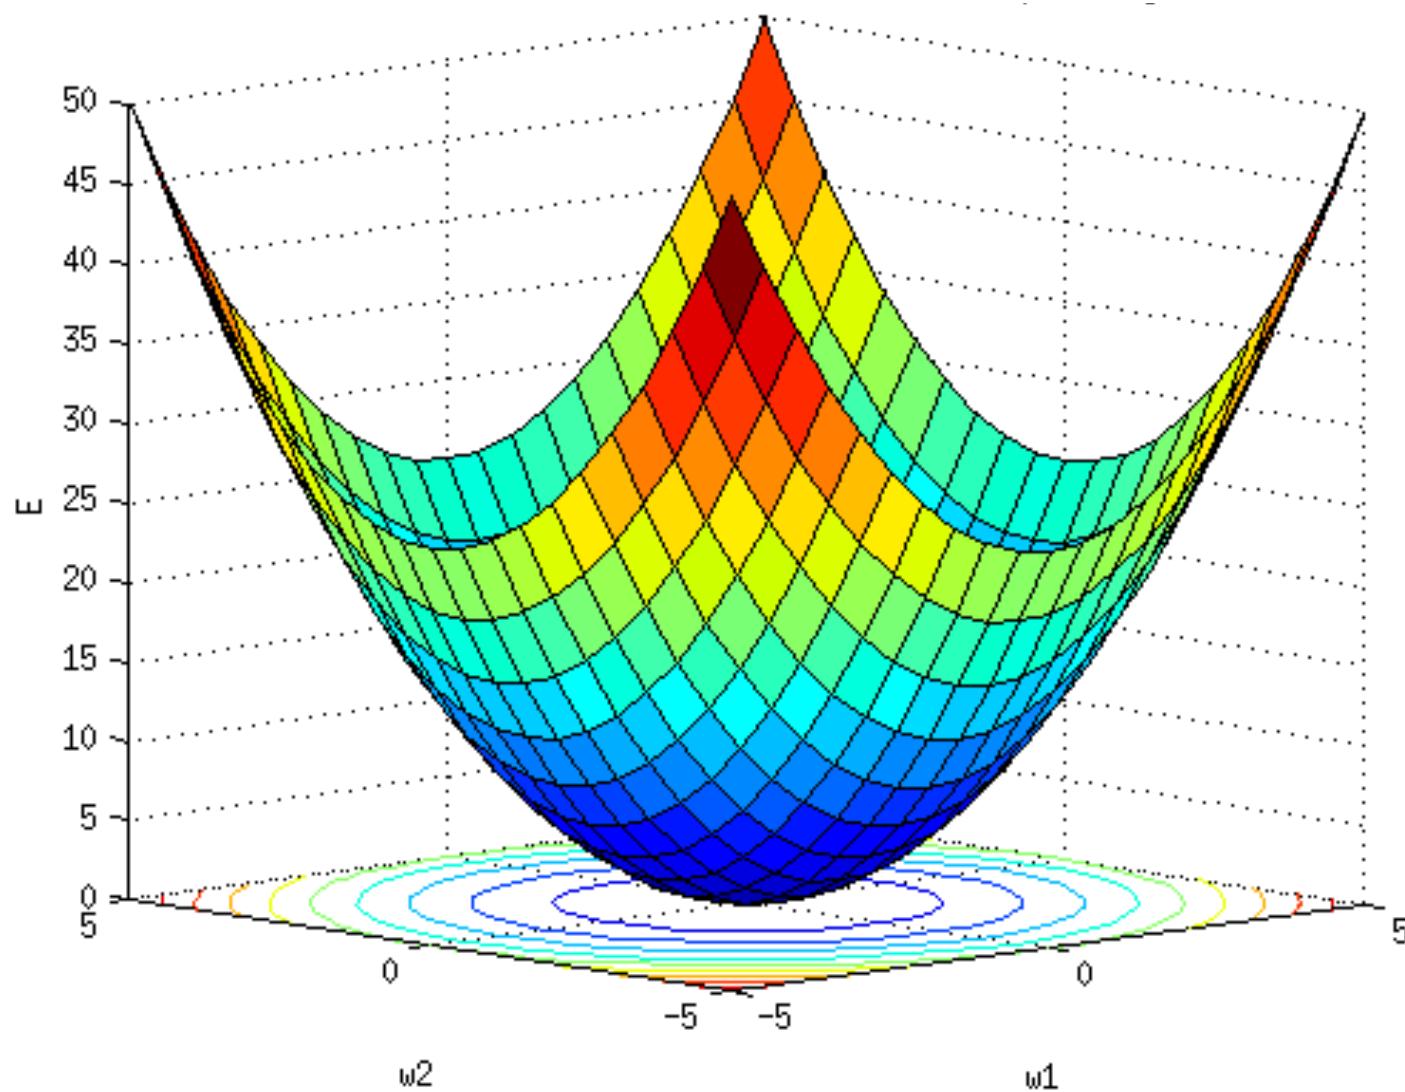
Backpropagation

Input



More details: https://www.youtube.com/watch?v=x_Eamf8MHwU

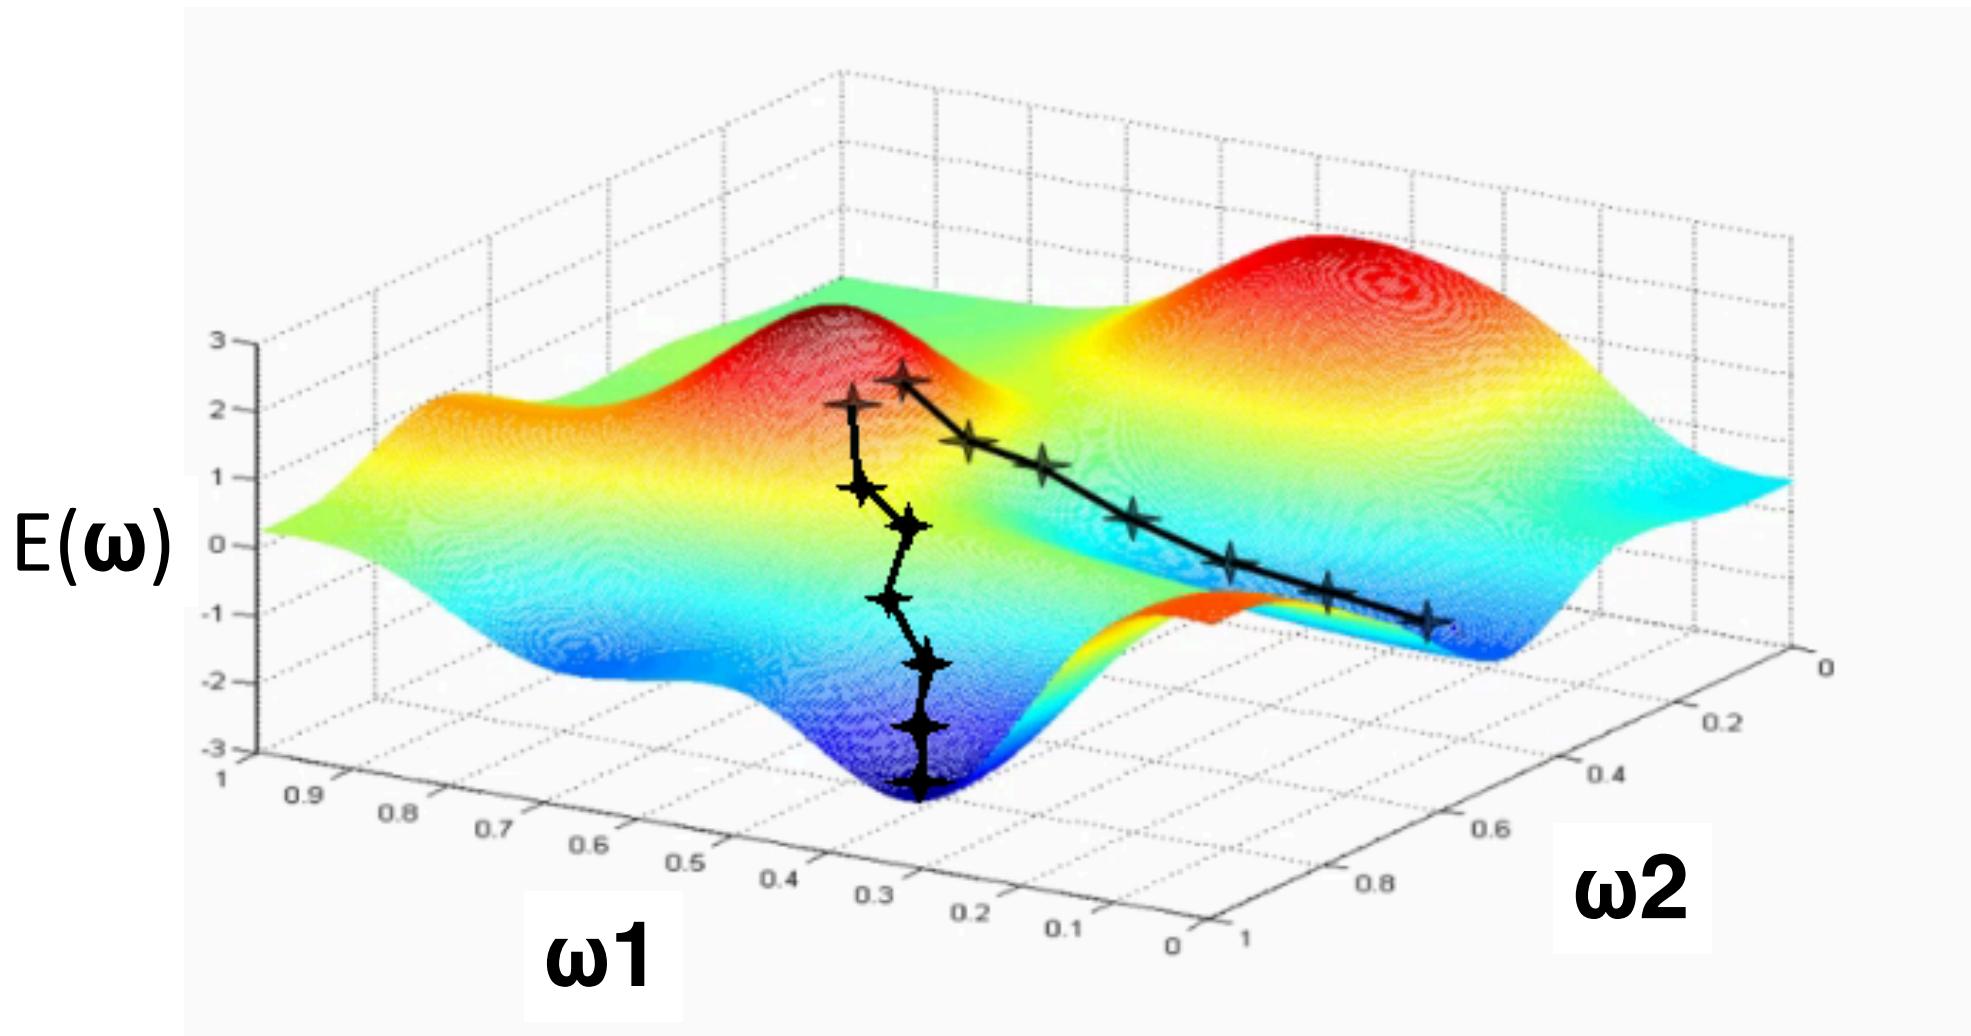
Ideal Cost Function



Real-world Cost Function



Gradient descend



More details: <https://www.youtube.com/watch?v=TveJaJs5Fqk>

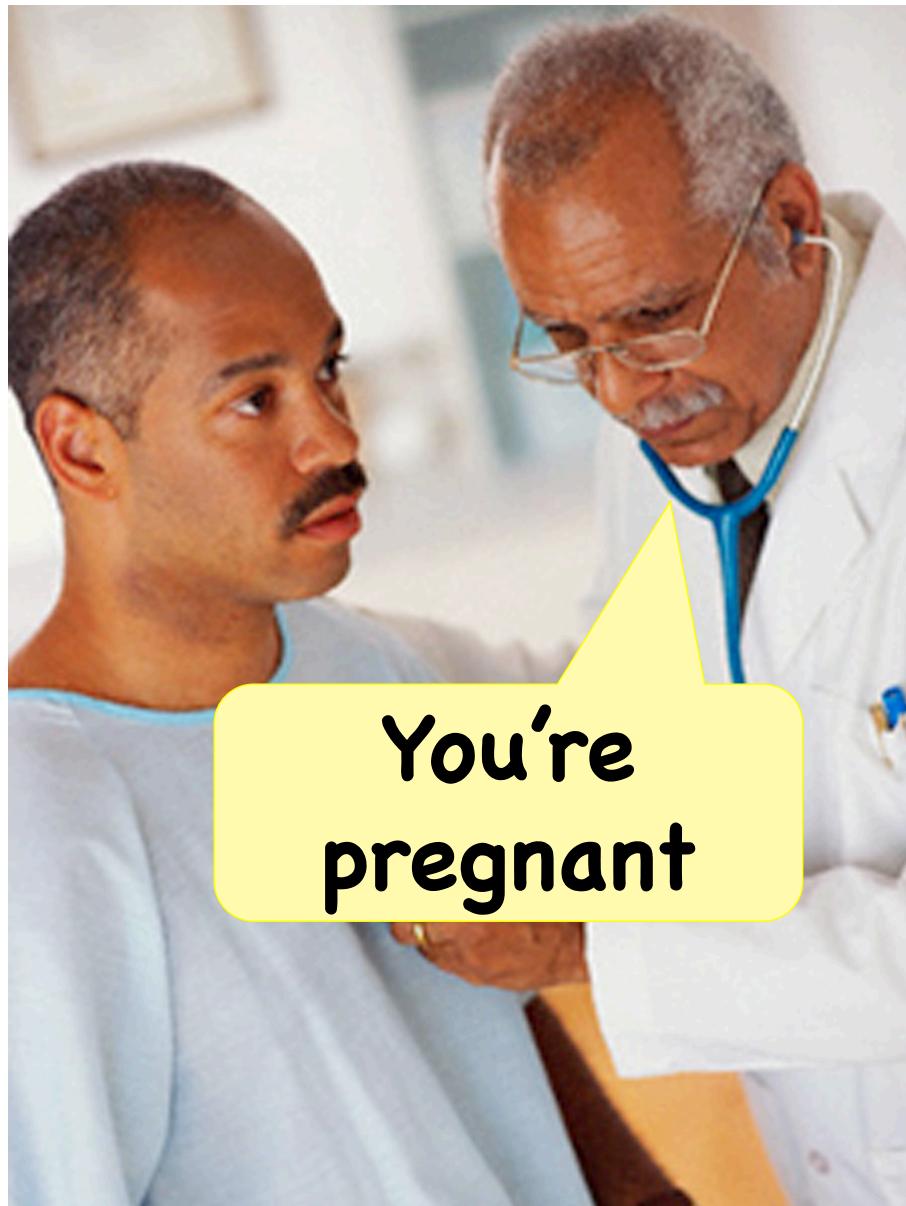
Iterate many times



Evaluate performance

- Confusion matrix
- Precision & Recall
- ROC curve (Receiver operating characteristic)

False Positive



False Negative



Confusion matrix

True class

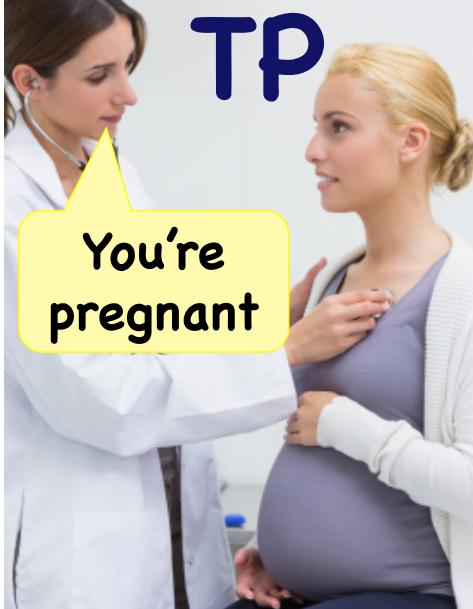
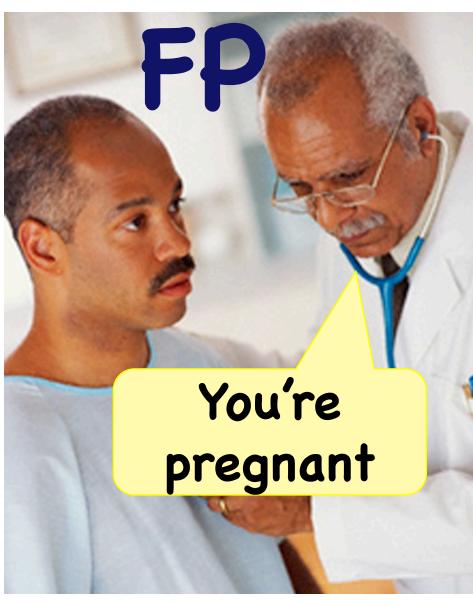
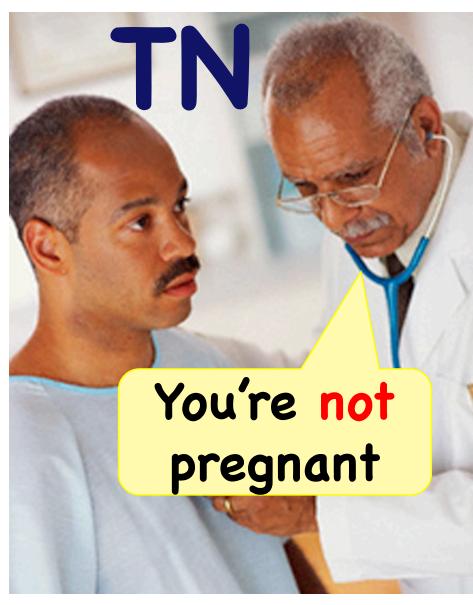
P

N

Predicted class

P

N

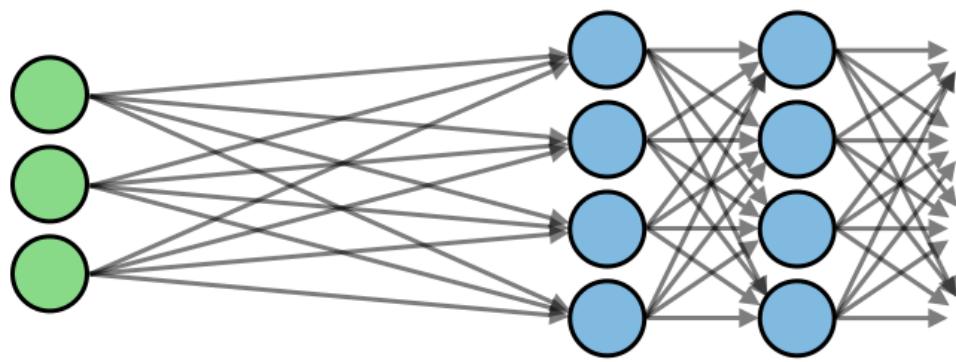
Some useful resources

- <https://playground.tensorflow.org>
- <https://seat.massey.ac.nz/personal/s.r.marsland/MLBook.html>
- <http://neuralnetworksanddeeplearning.com/>
- [Python machine learning](#) - Sebastian Raschka
- [Deep learning with Python](#) - Francois Chollet
- [Machine Learning - An Algorithmic Perspective](#) - Stephen Marsland
- Datasets put together by Men-Andrin and Zach Ross
 - <http://scedc.caltech.edu/research-tools/deeplearning.html>

Many pictures in this talk are from internet, I thank all the authors here!

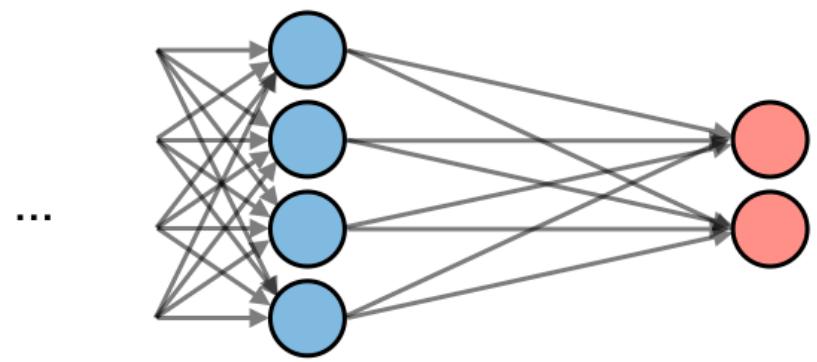
Go to the notebook

Deep neural network



Input layer

Hidden layer 1



...

... Hidden layer k

Output layer

Unreasonable Benefits of Deep Learning

simplicity

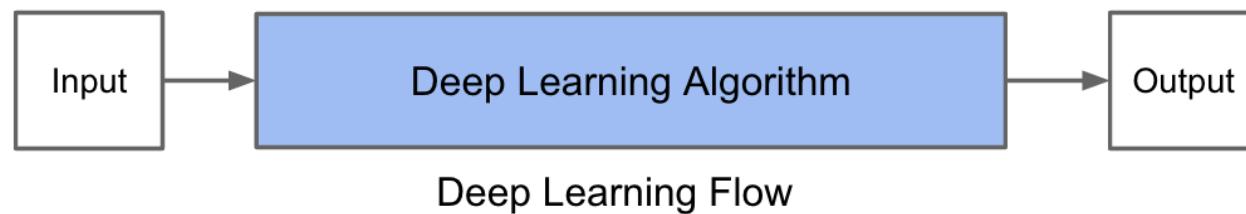
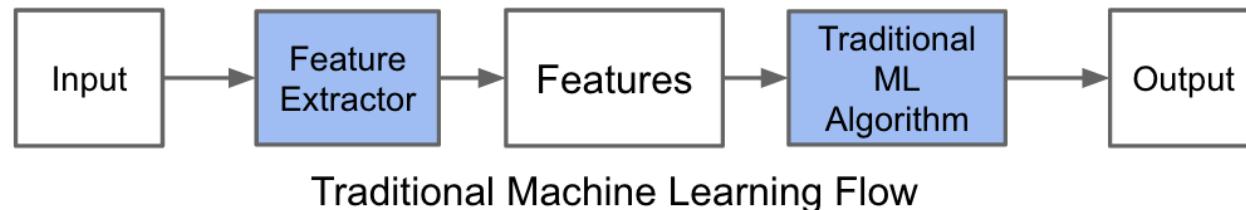
accuracy

flexibility

hacks

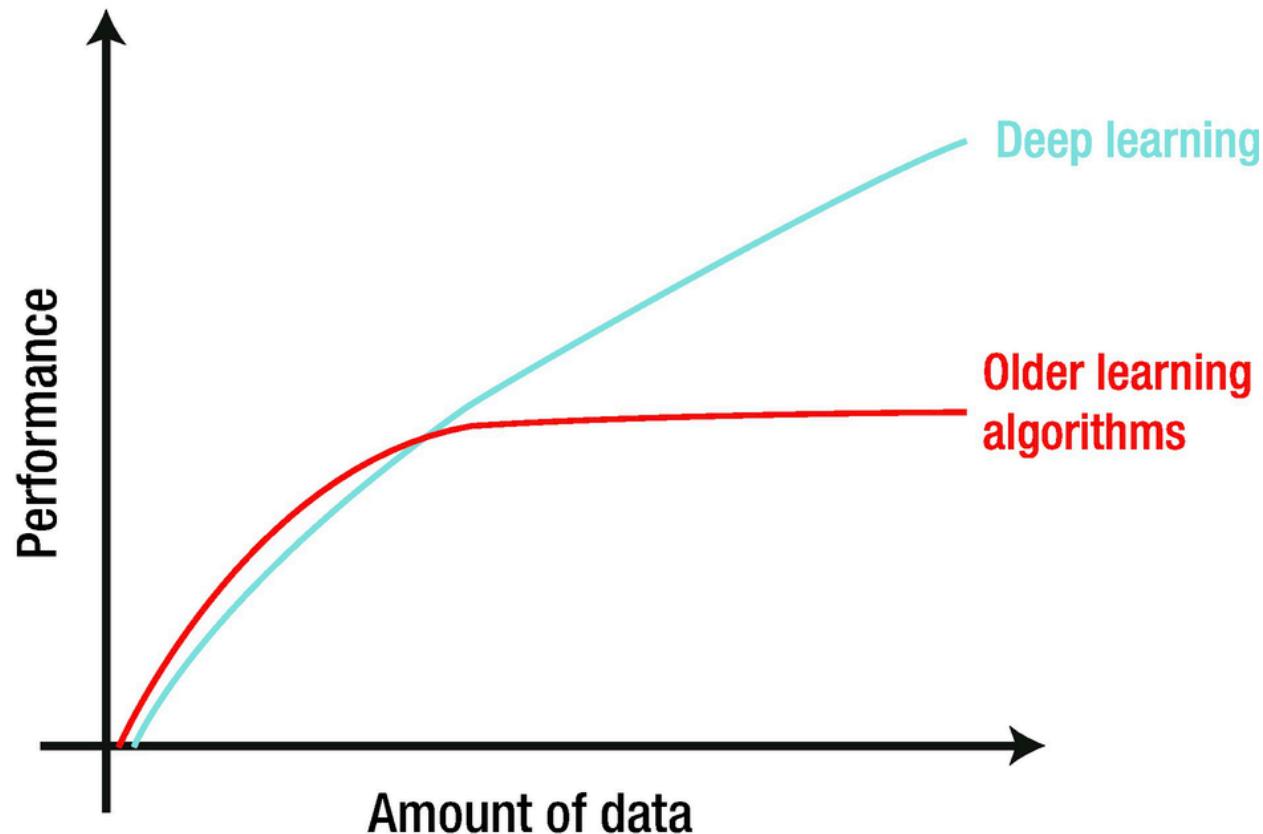
Benefits of deep learning

Simplicity



Benefits of deep learning

Accuracy

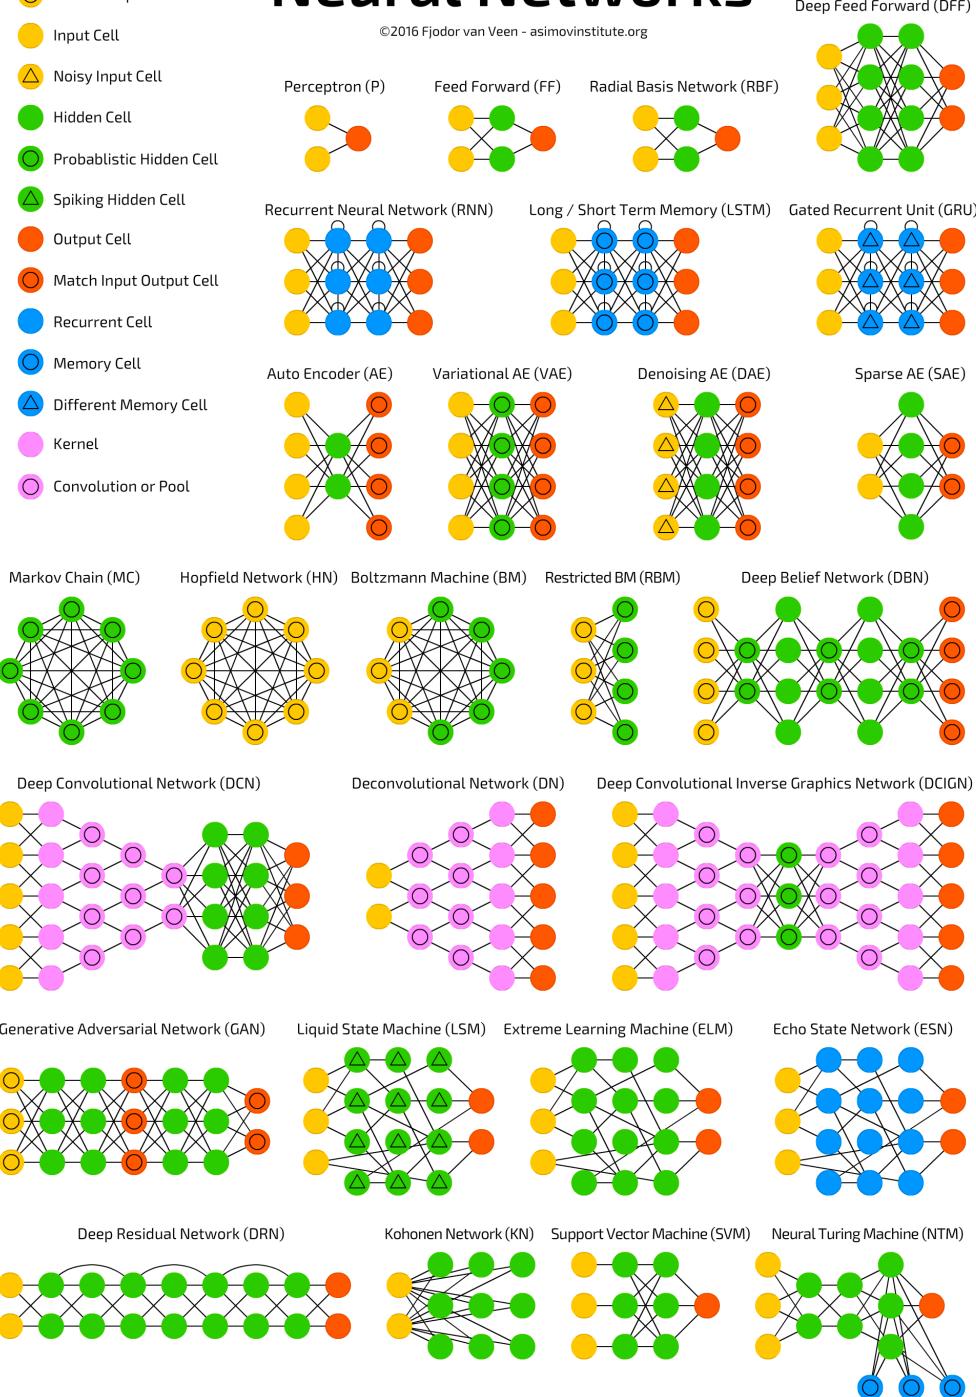


A mostly complete chart of

Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org

- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

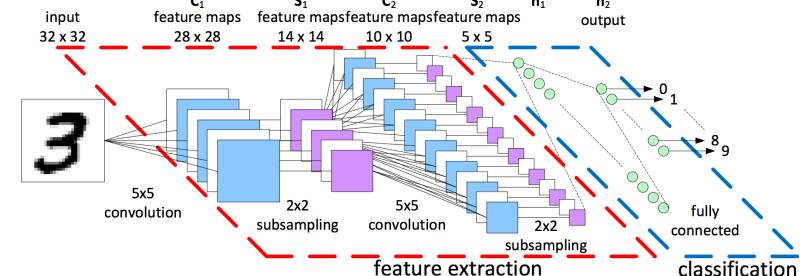


<https://towardsdatascience.com/the-mostly-complete-chart-of-neural-networks-explained-3fb6f2367464>

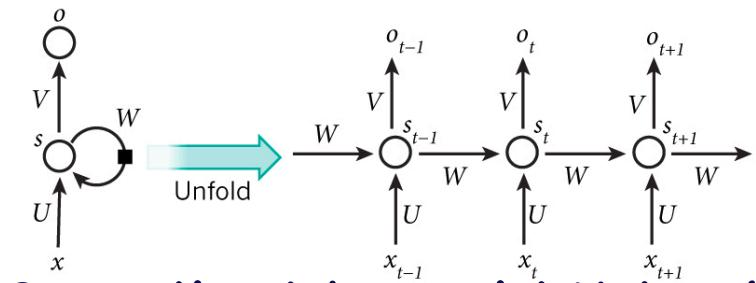
Benefits

Flexible

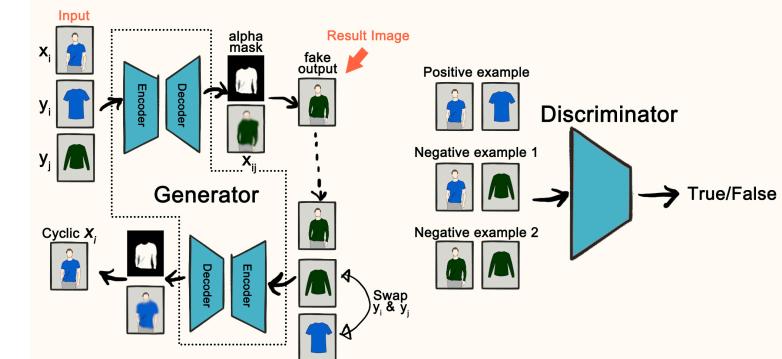
Convolutional Neural Network



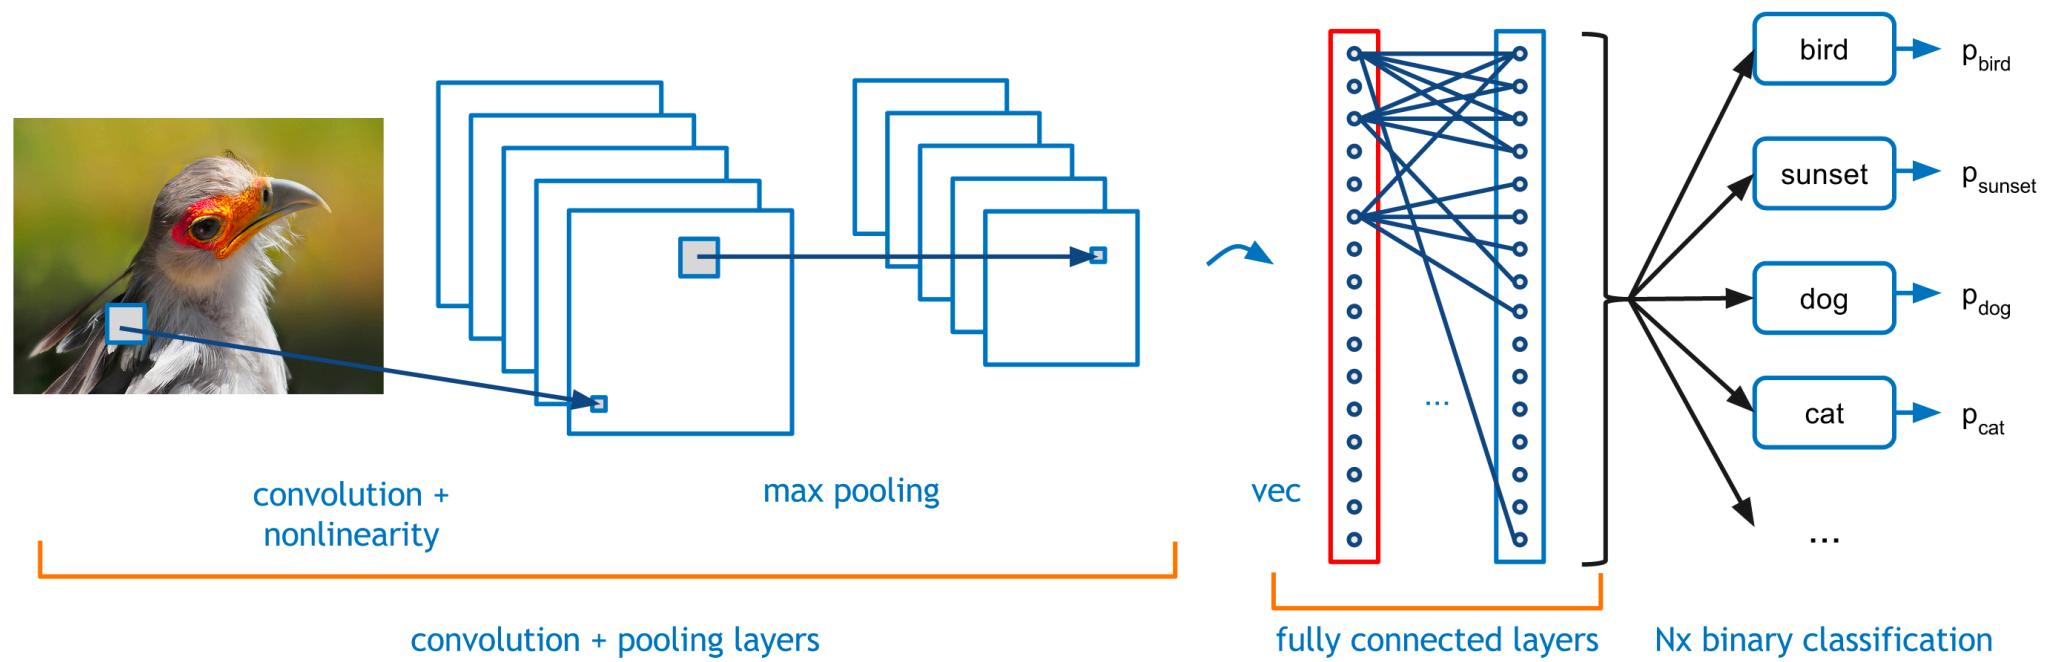
Recurrent Neural Network



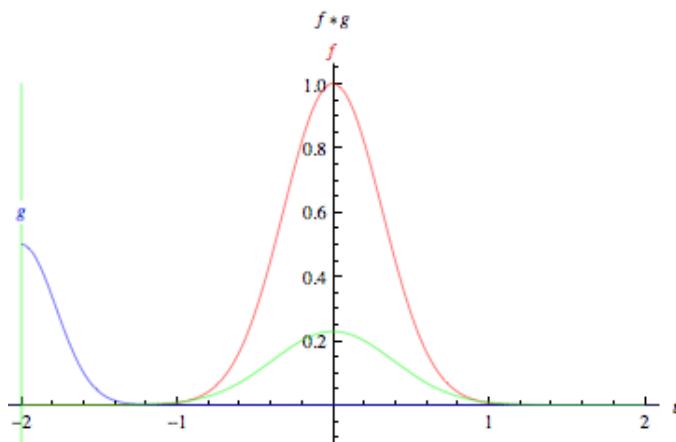
Generative Adversarial Network



Convolutional neural network



Convolution operation



<http://mathworld.wolfram.com/>

1 x1	1 x0	1 x1	0	0
0 x0	1 x1	1 x0	1	0
0 x1	0 x0	1 x1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Max pooling

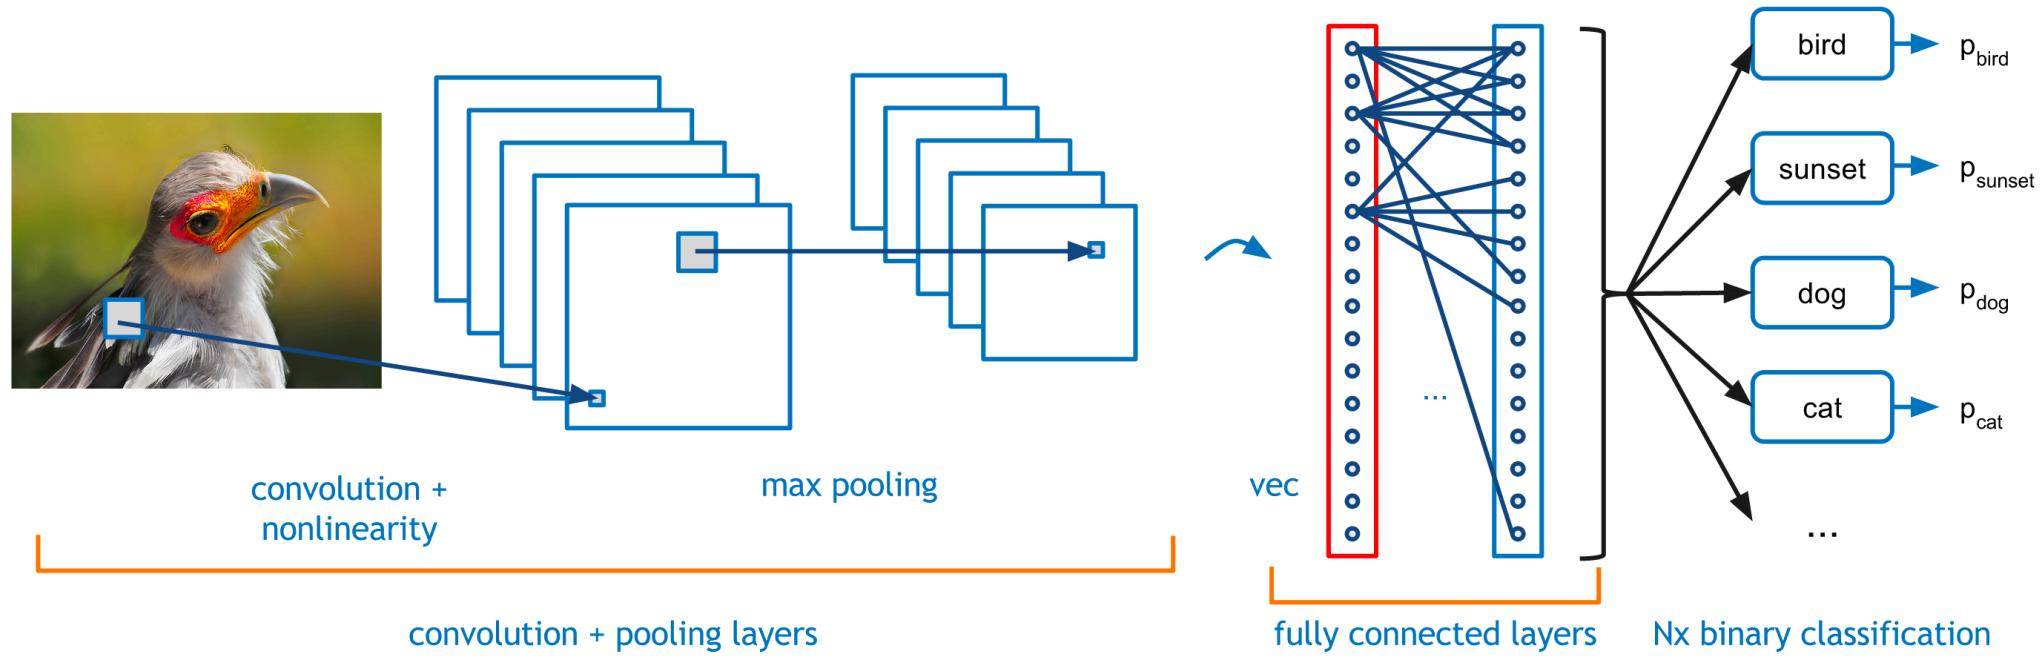
1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2
window and stride 2

The diagram illustrates the process of max pooling. On the left, a 4x4 input matrix is shown with values: Row 1: 1, 1, 2, 4; Row 2: 5, 6, 7, 8; Row 3: 3, 2, 1, 0; Row 4: 1, 2, 3, 4. The matrix is color-coded in a checkerboard pattern. A horizontal arrow points from the input to the output on the right. The output is a 2x2 matrix with values: Top-left cell (row 1, column 1) is 6 (from the max of 1, 1); Top-right cell (row 1, column 2) is 8 (from the max of 2, 4); Bottom-left cell (row 2, column 1) is 3 (from the max of 5, 6); Bottom-right cell (row 2, column 2) is 4 (from the max of 7, 8). The output matrix is also color-coded.

6	8
3	4

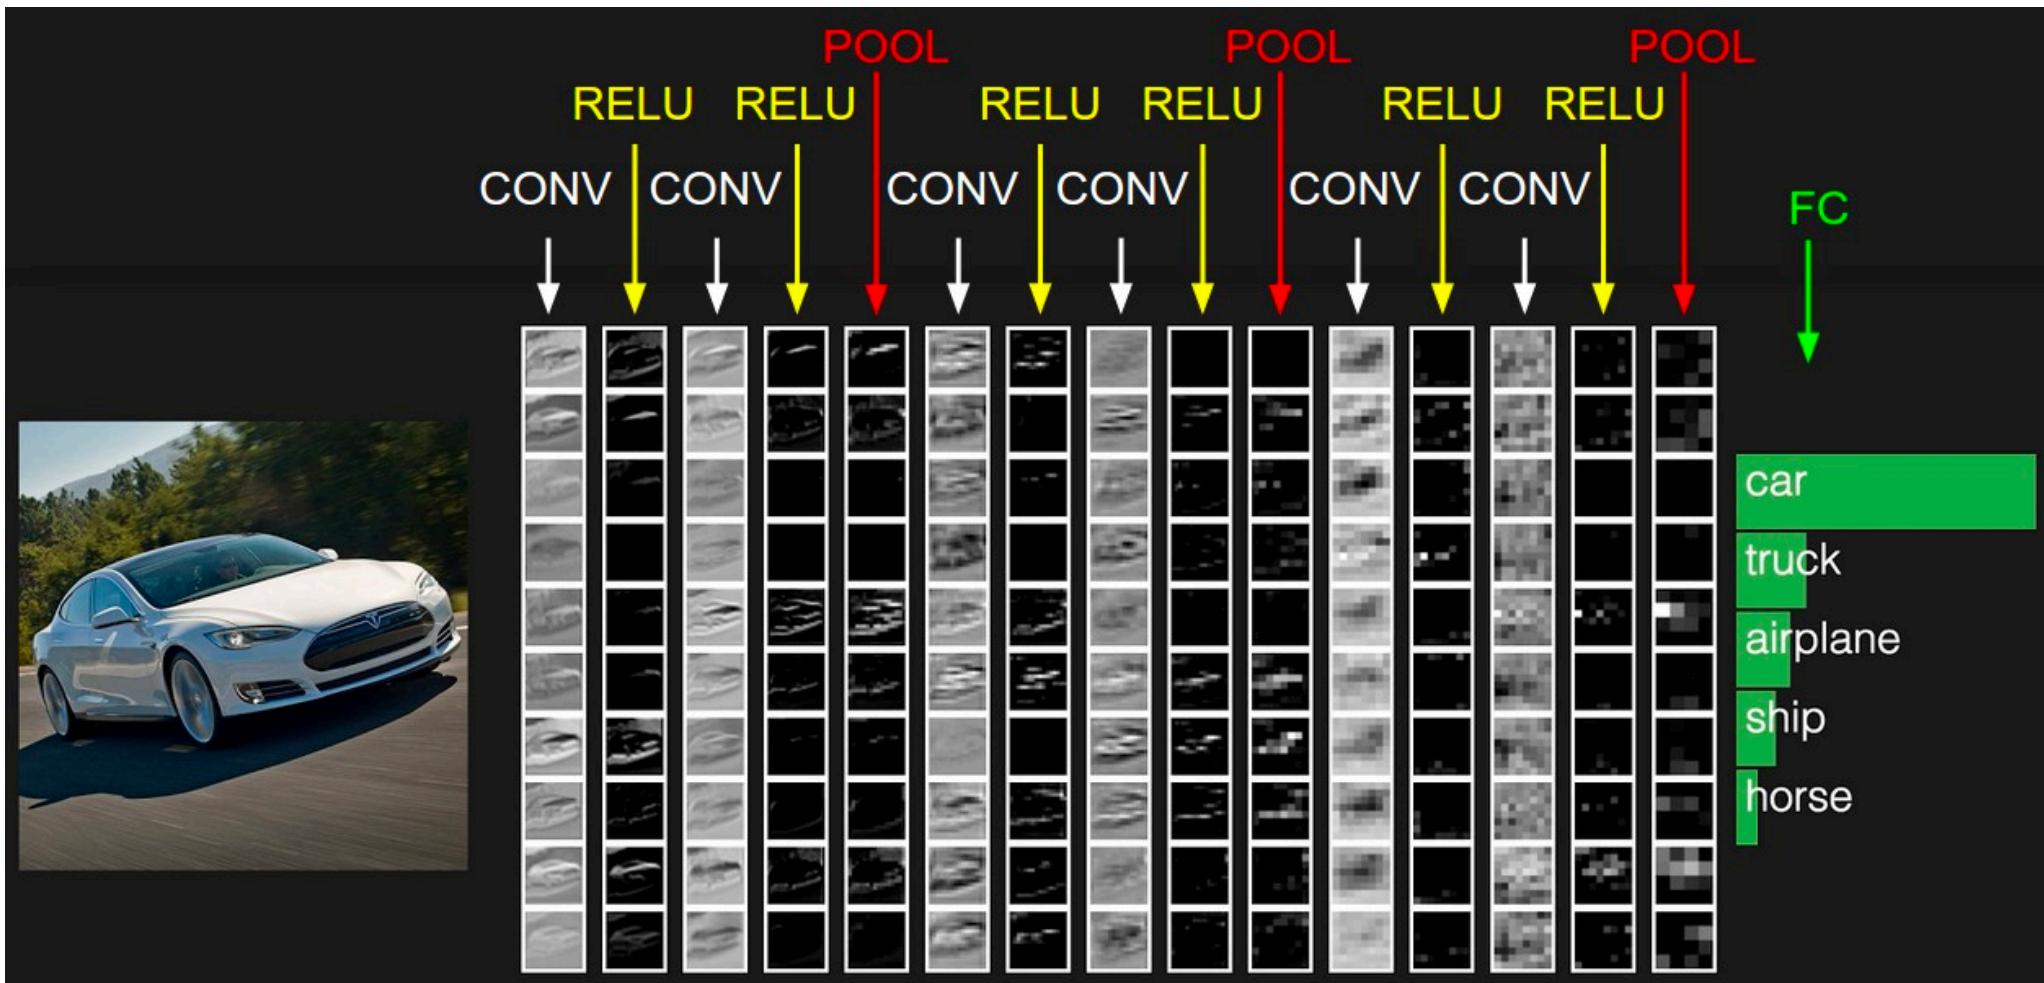
Convolutional neural network



Feature extraction

Classification

CNN example



Some useful resources

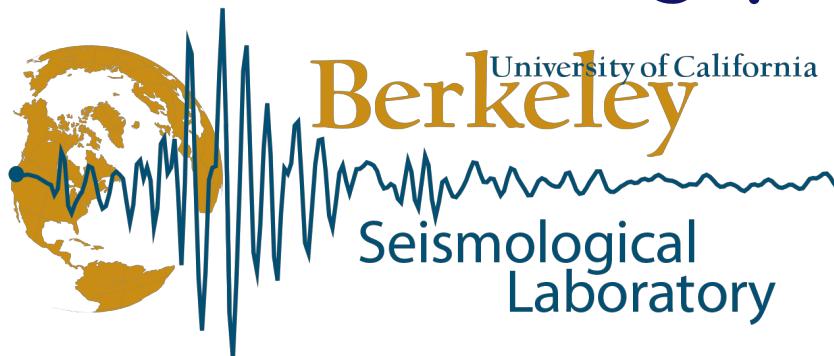
- <http://cs231n.github.io/convolutional-networks/>
- <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>
- <http://neuralnetworksanddeeplearning.com/>
- [Deep learning with Python](#) - Francois Chollet
- [Deep learning](#) - Ian Goodfellow
- Datasets put together by Men-Andrin and Zach Ross
 - <http://scedc.caltech.edu/research-tools/deeplearning.html>

Many pictures in this talk are from internet, I thank all the authors here!

Go to the notebook



Thanks
kongqk@Berkeley.edu



<http://seismo.berkeley.edu/qingkaikong/>